

# RecogniChess

An unsupervised domain-adaptation approach to chessboard recognition

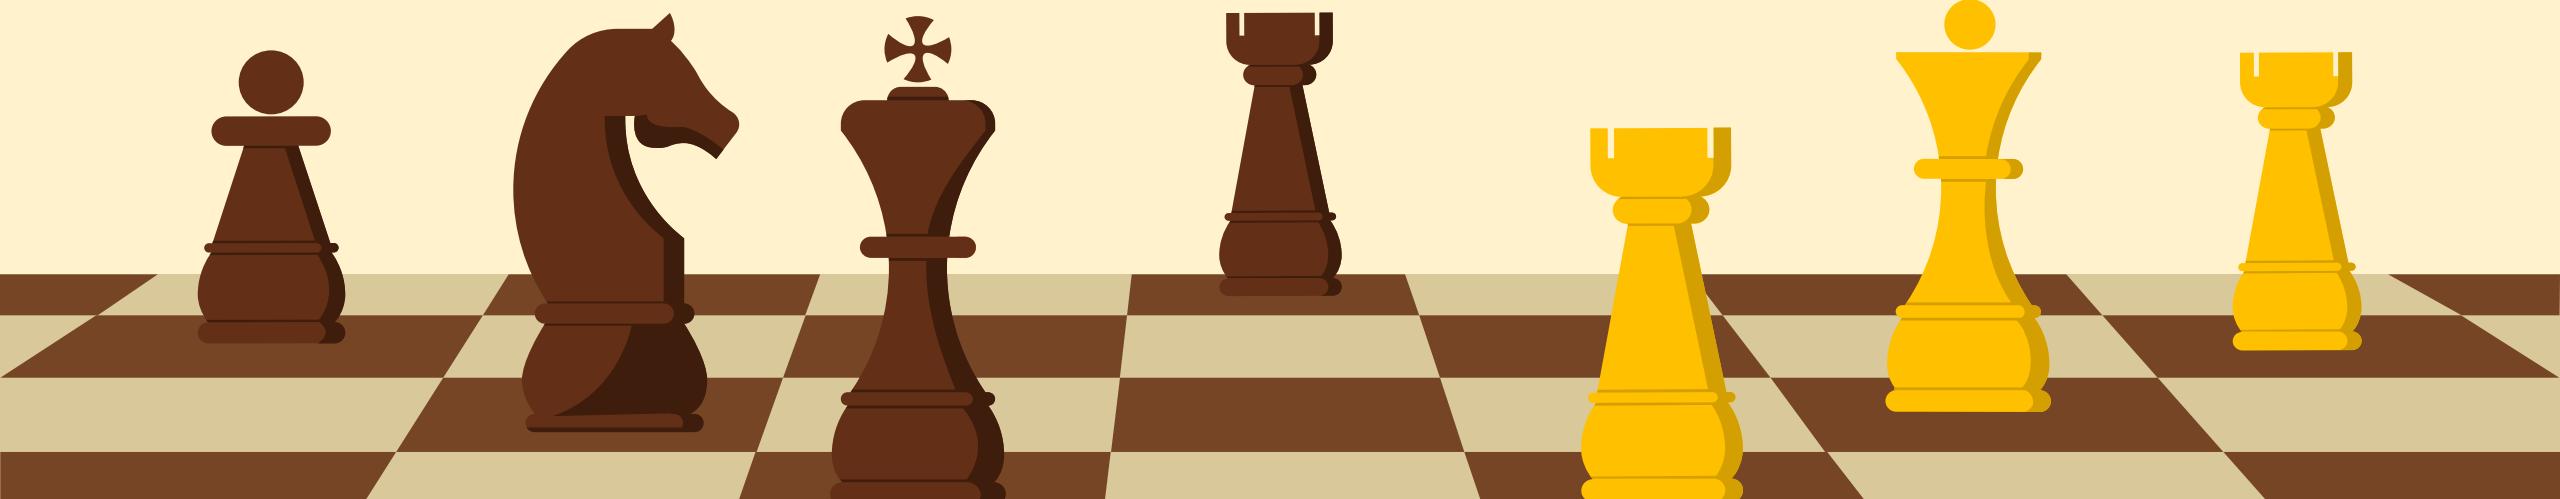
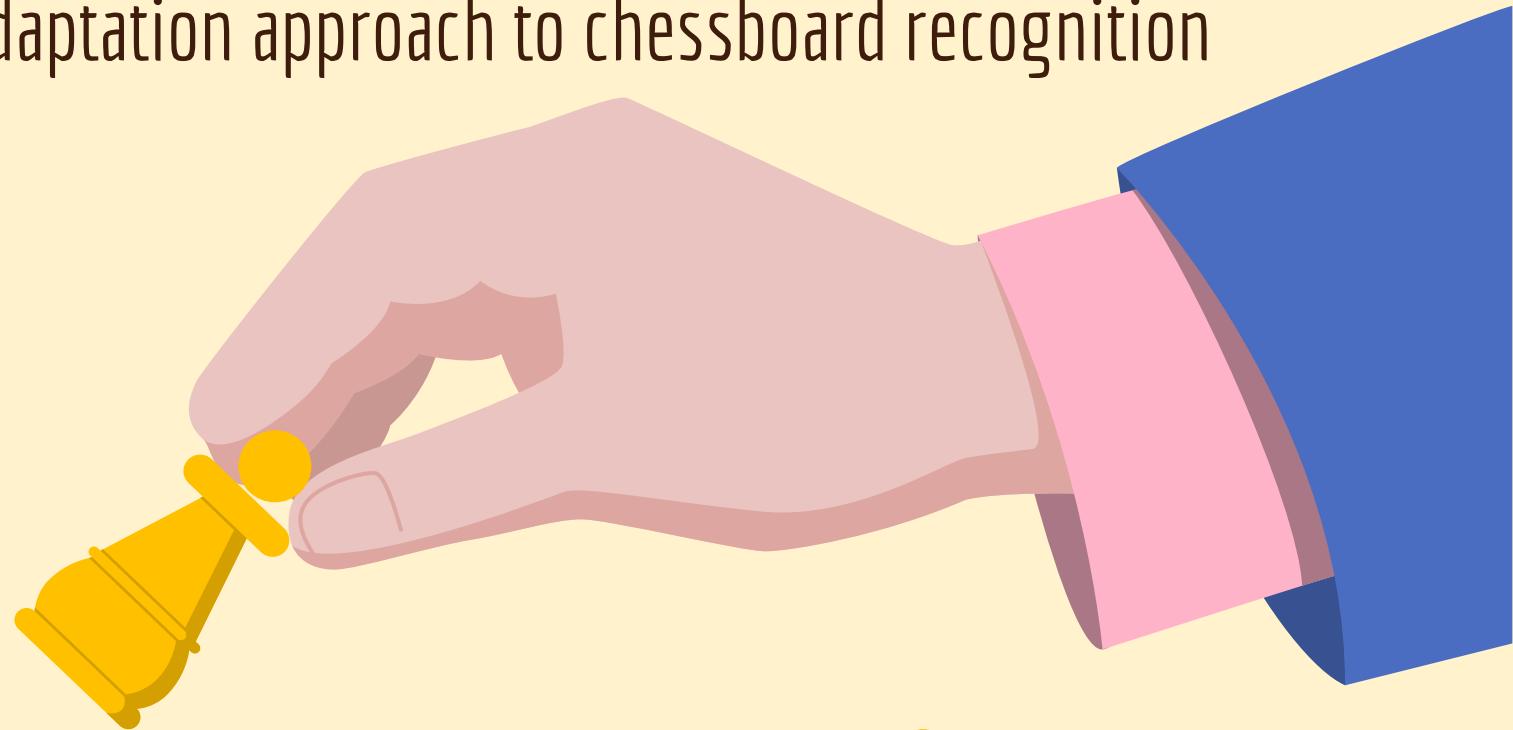
The team:

Wassim Jabbour

Enzo Benoit-Jeannin

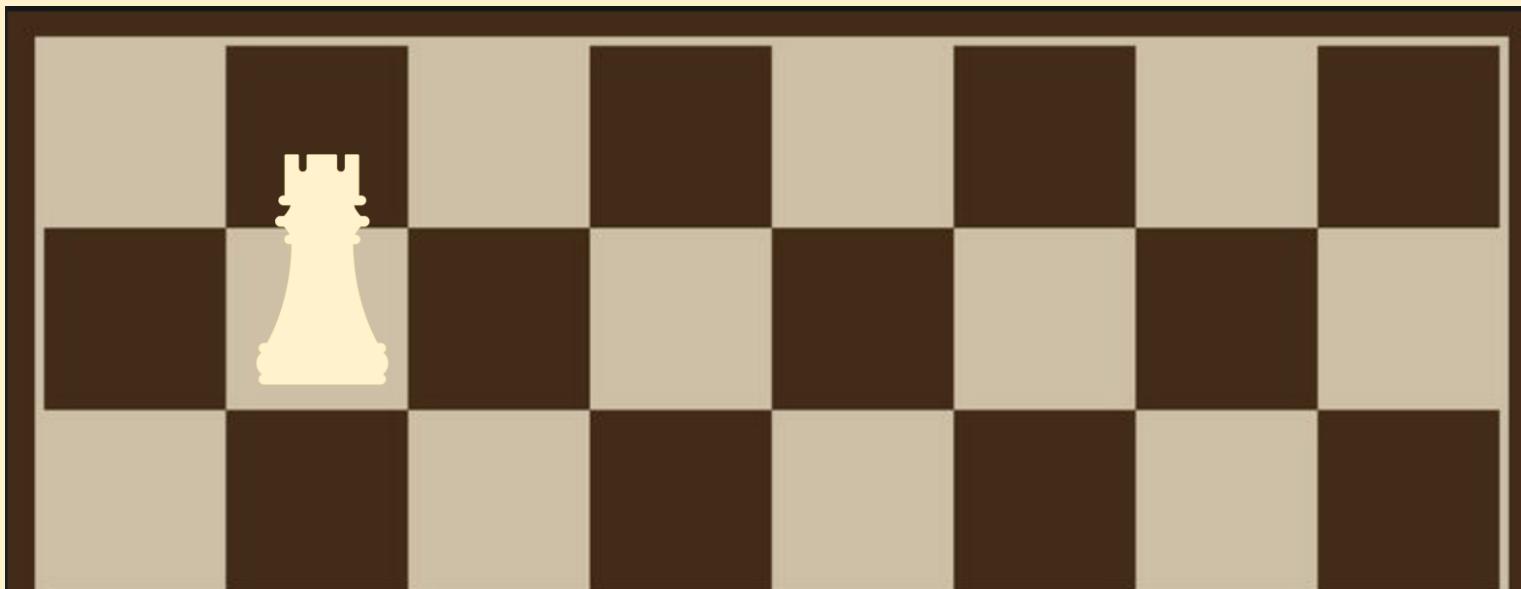
Oscar Bedford

Saif Shahin



# Introduction

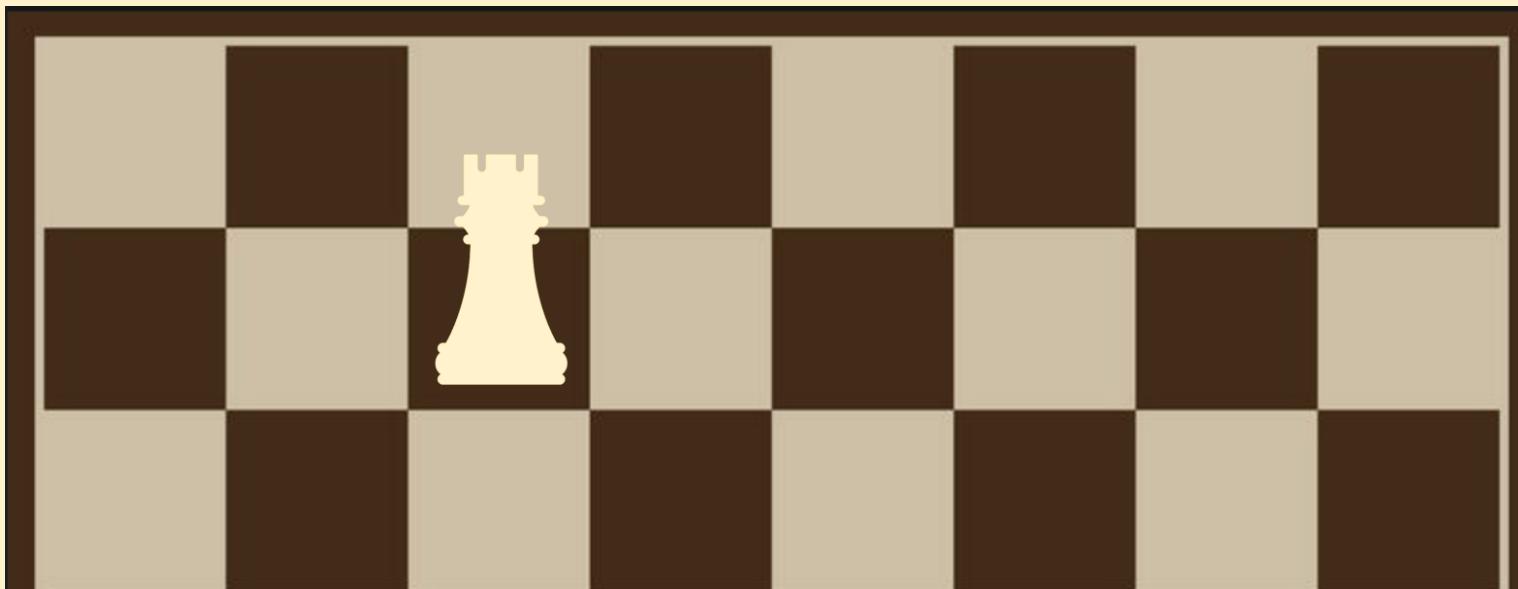
Chess move  
annotation is  
important



# Introduction

Chess move  
annotation is  
important

Automation  
can be highly  
beneficial

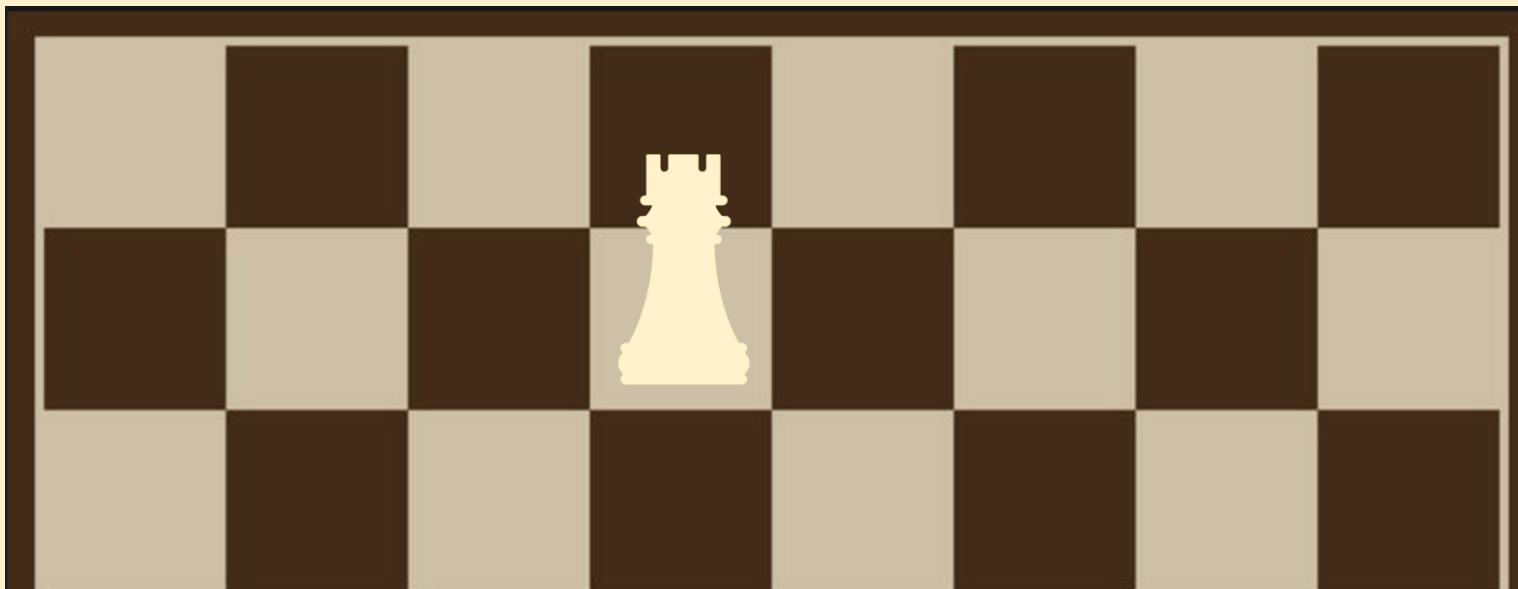


# Introduction

Chess move  
annotation is  
important

Automation  
can be highly  
beneficial

Requires a lot  
of labeled data



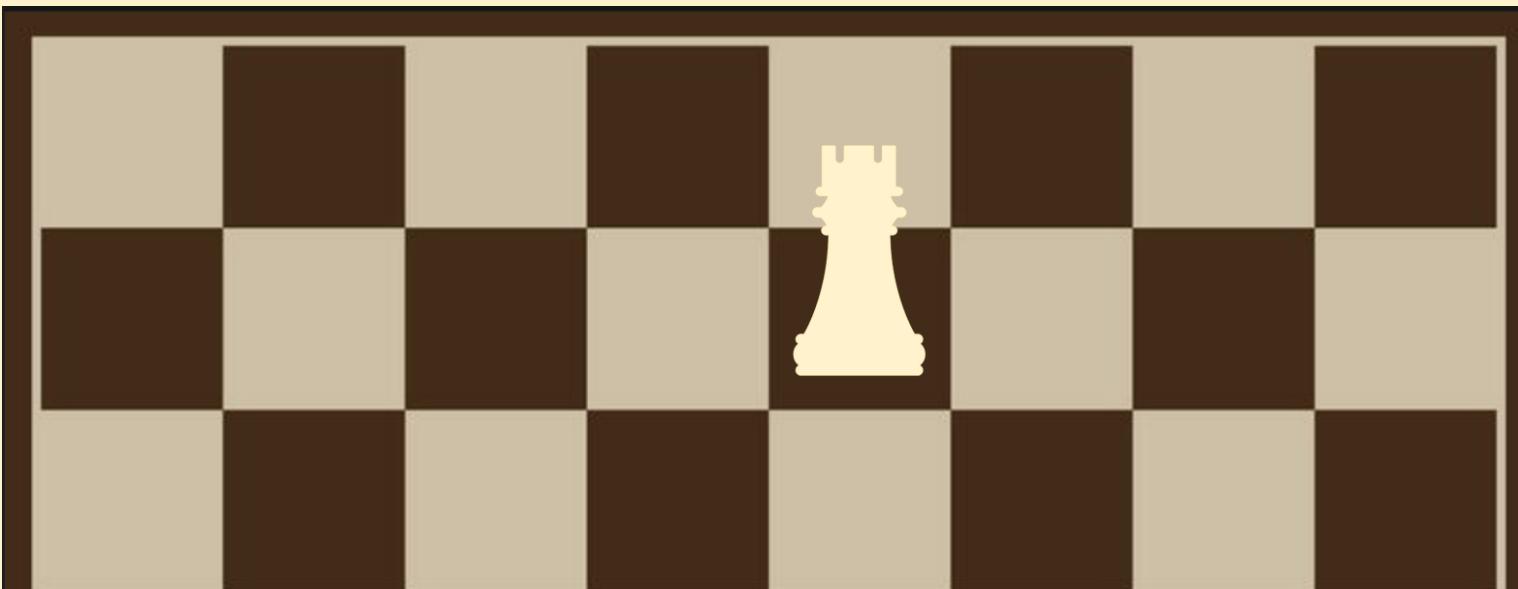
# Introduction

Chess move  
annotation is  
important

Automation  
can be highly  
beneficial

Requires a lot  
of labeled data

Labeled data  
is hard to  
come by



# Introduction

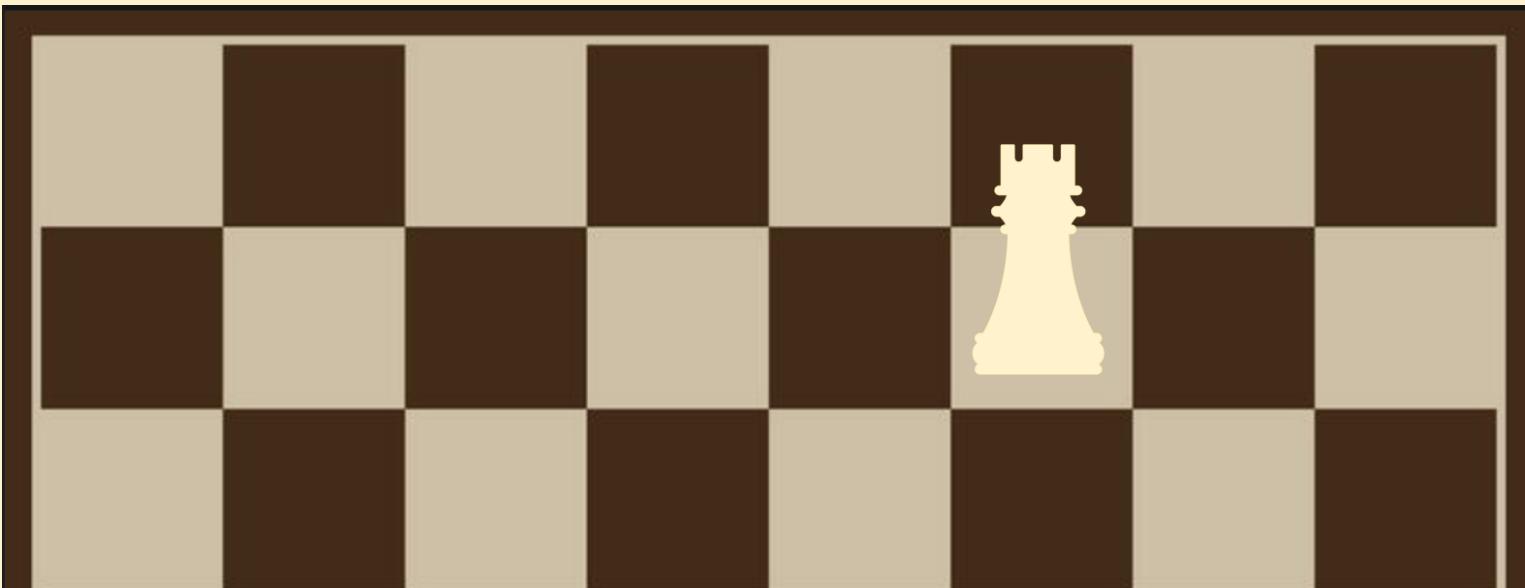
Chess move  
annotation is  
important

Automation  
can be highly  
beneficial

Requires a lot  
of labeled data

Labeled data  
is hard to  
come by

Labelling  
chess board  
images is  
tedious



# Introduction

Chess move annotation is important

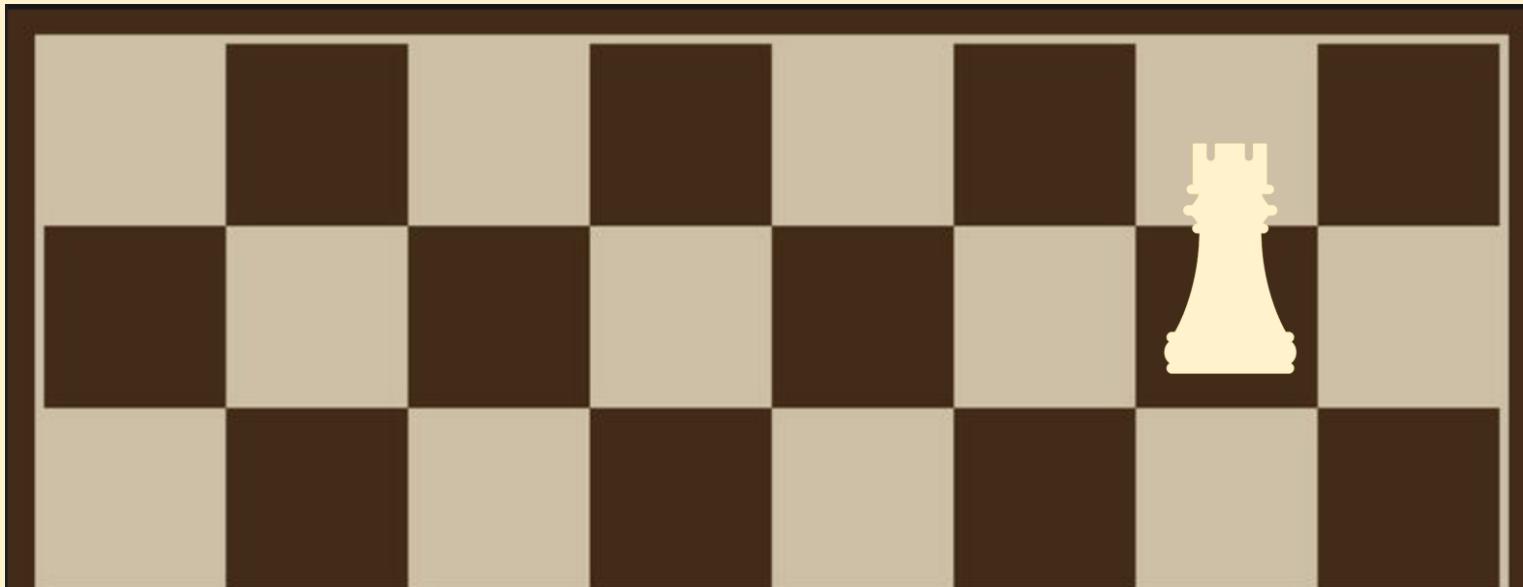
Automation can be highly beneficial

Requires a lot of labeled data

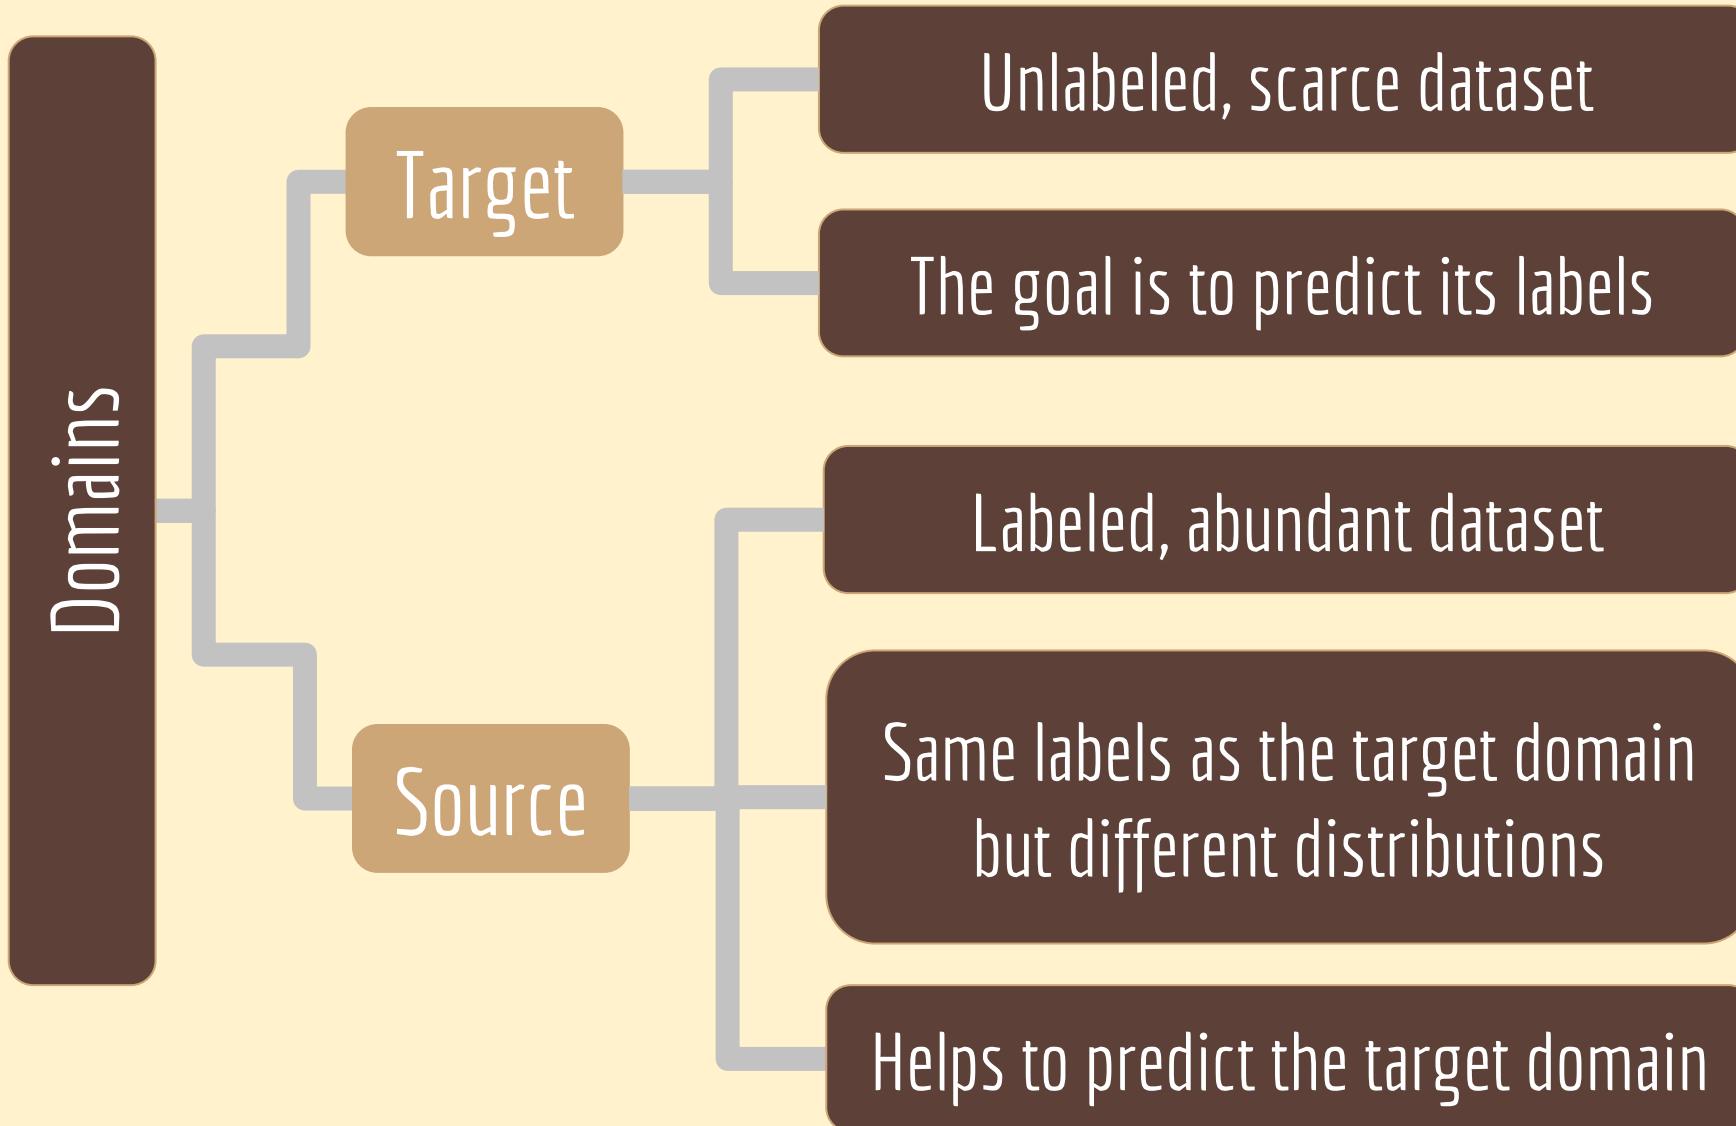
Labeled data is hard to come by

Labelling chess board images is tedious

Unsupervised domain adaptation



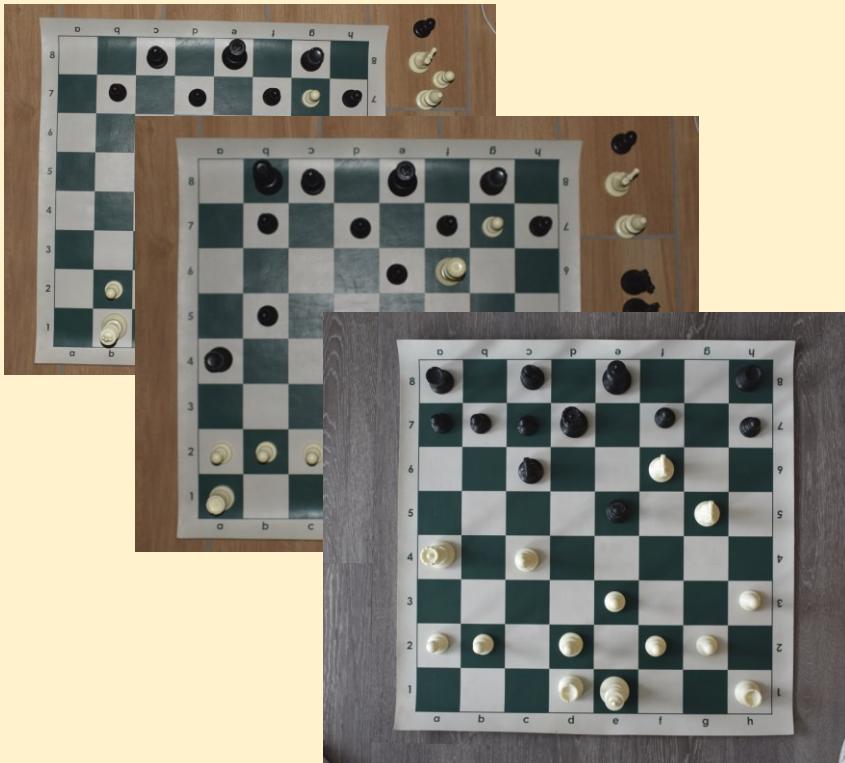
# Overview of unsupervised DA<sup>[1]</sup>



[1] S. Ben-David et al., "A theory of learning from different domains," *Mach. Learn.*, vol. 79, pp. 151-175, 2010.

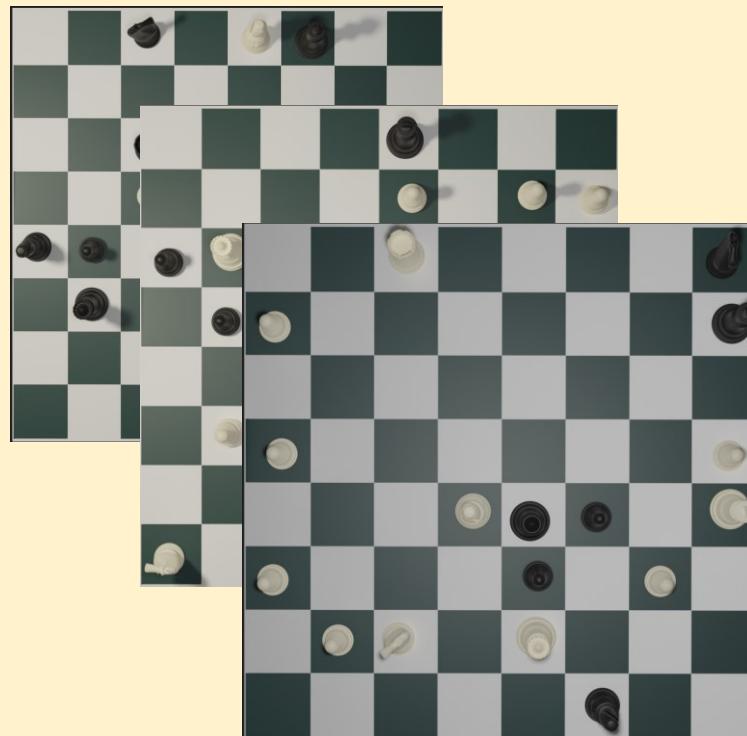
# Unsupervised DA in the context of chess

Capture unprocessed,  
unlabeled, top-view  
chessboard images



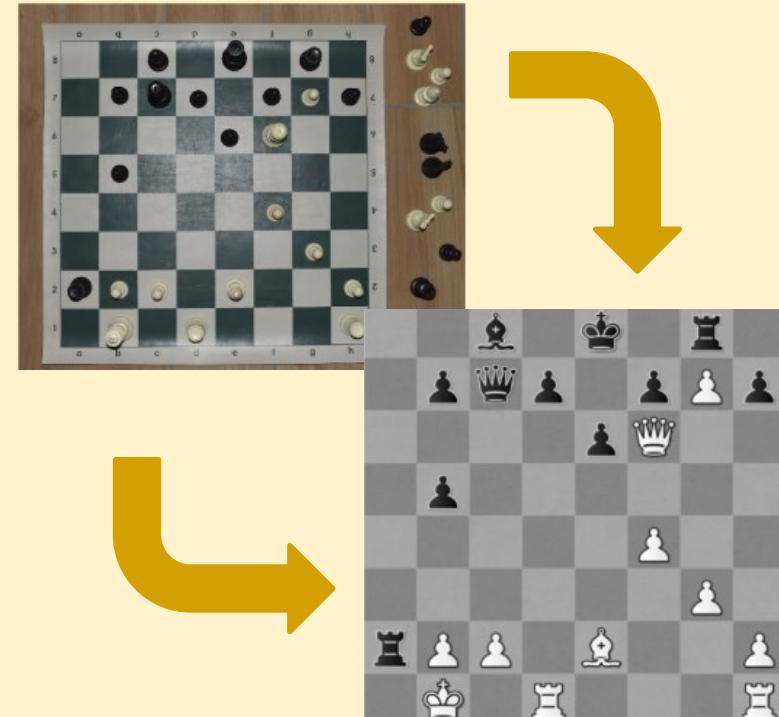
Target domain

Generate labeled images  
with similar feature / target  
distributions



Source domain

Build a pipeline that utilizes  
the source domain inputs to  
predict the target domain



Using Unsupervised DA

# Target domain dataset

500 full, distinct, legal chess positions

Unlabelled (for the most part)

Different lighting and background conditions



[2] A. D. S. D. Neto and R. M. Campello, "Chess position identification using pieces classification based on synthetic images generation and deep neural network fine-tuning," in Proc. 21st Symp. Virtual Augmented Reality (SVR), pp. 152-160, IEEE, Oct. 2019.

# Source domain dataset

Generated using Blender

4500 distinct, labelled, and legal chess positions  
pre-processed into 288,000 individual squares

Different lighting conditions

Different centering and rotation of the pieces

Multiple top-view camera angles

Makes the trained model invariant to all of  
the above conditions

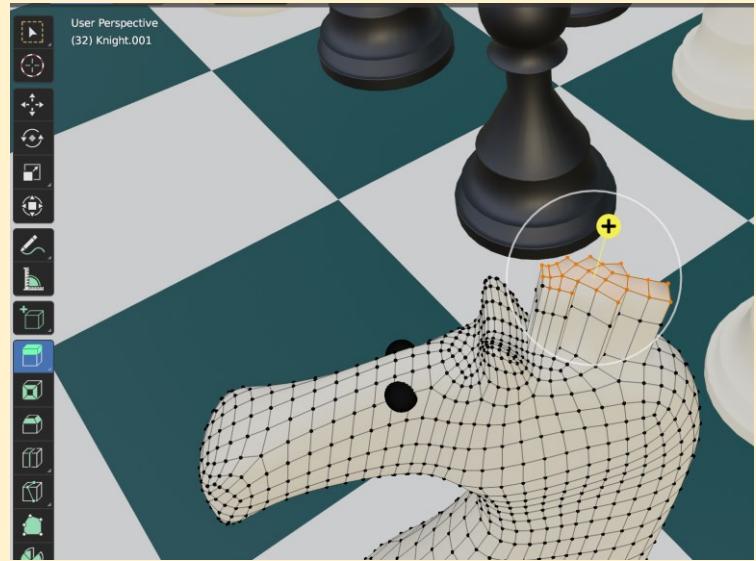
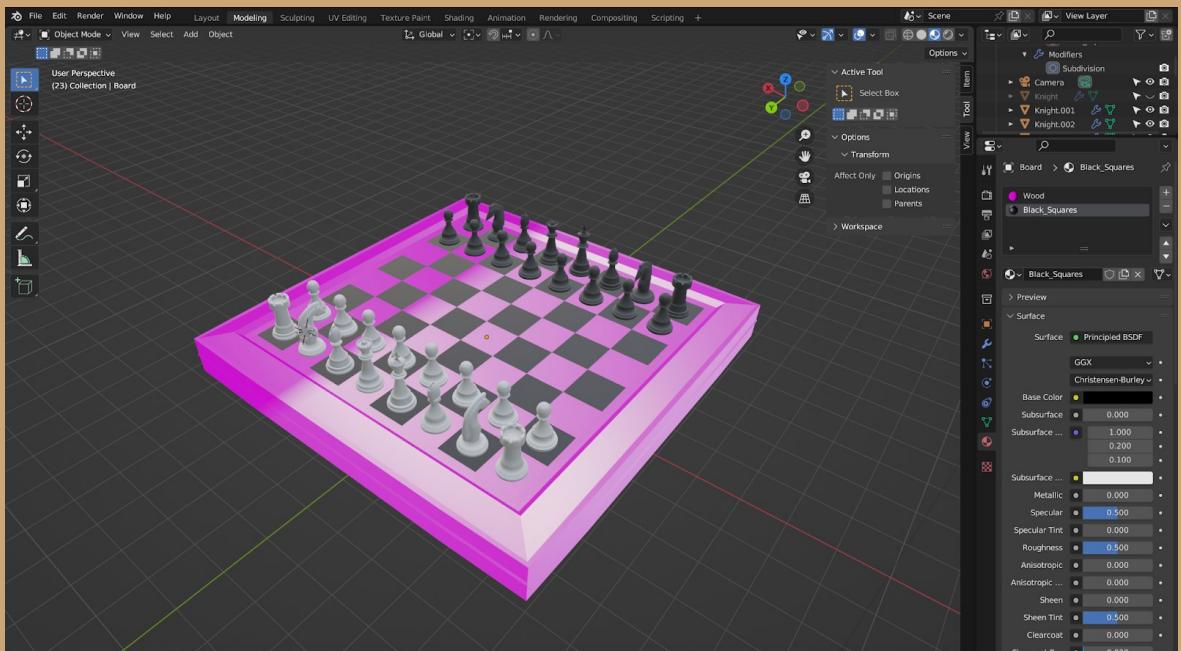


# Data generation details

Utilized Blender to generate top-view chessboard images

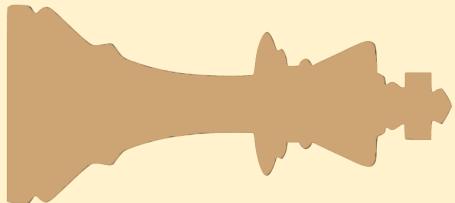
Tuned the texture and shapes of the pieces to match the target dataset to the best possible extent

Camera angle and lighting were varied between positions

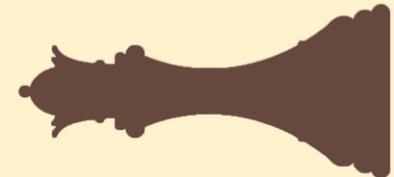


# Note about generating full images & cropping them

For the synthetic dataset, full positions were generated & cropped into squares



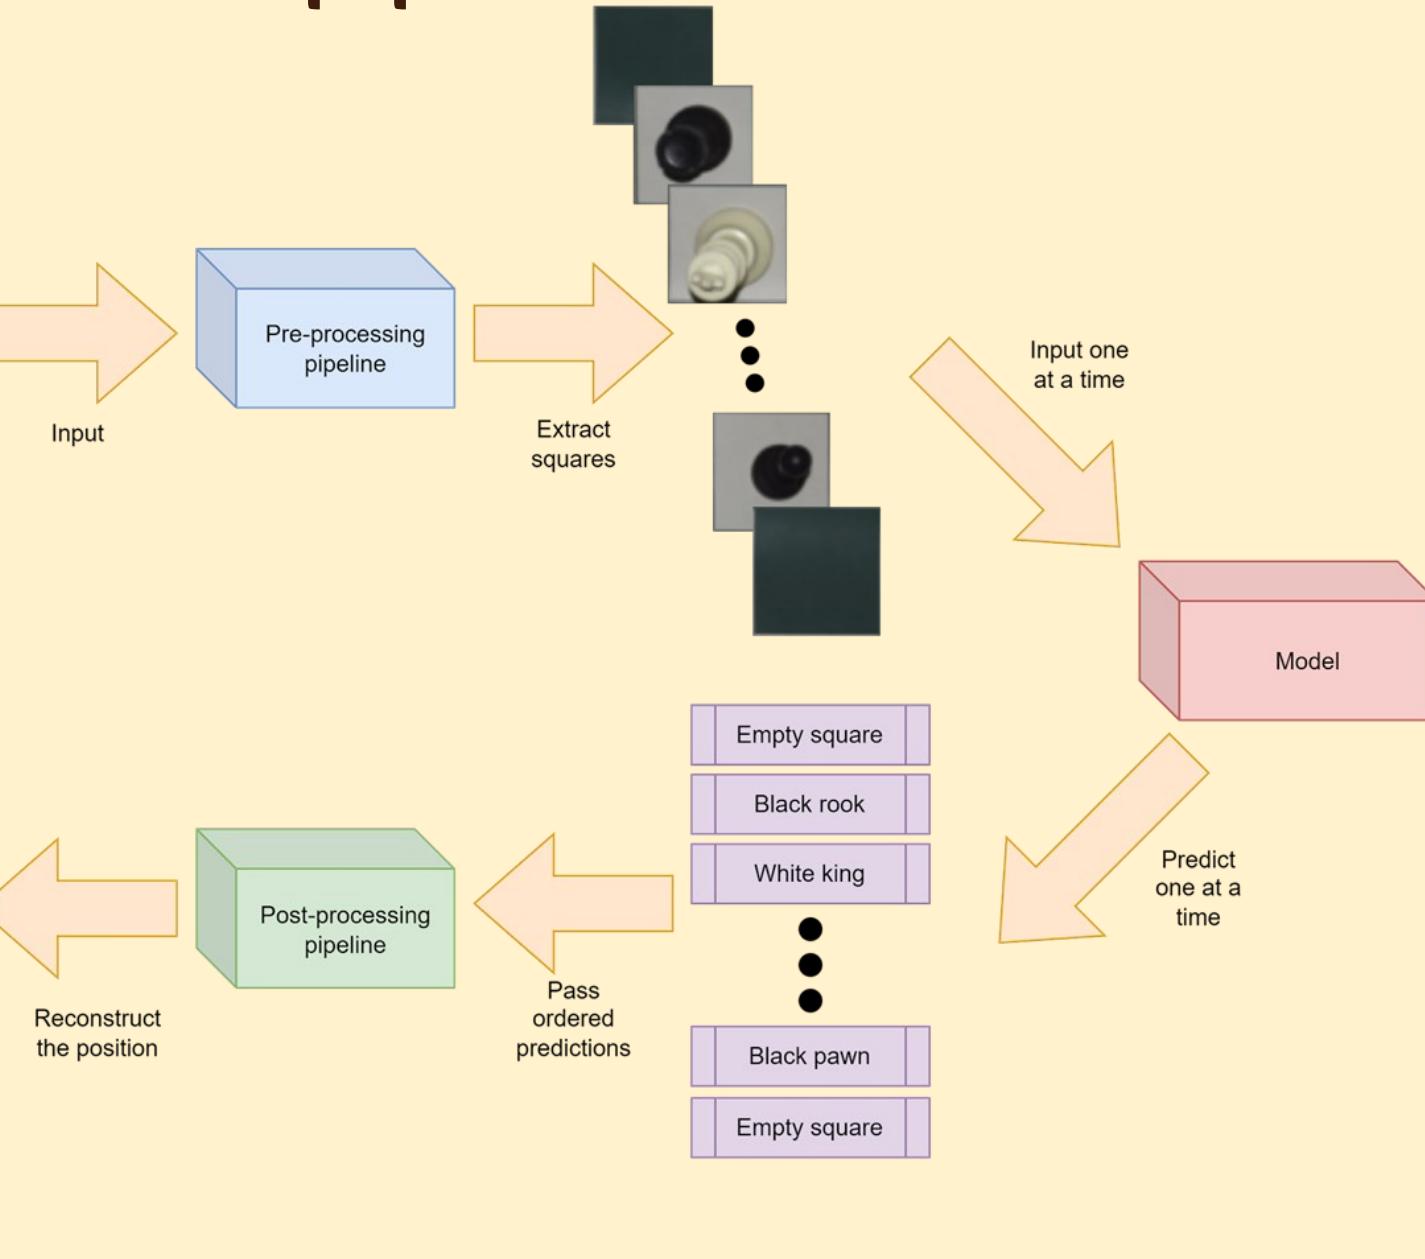
Why not generate individual squares directly ?



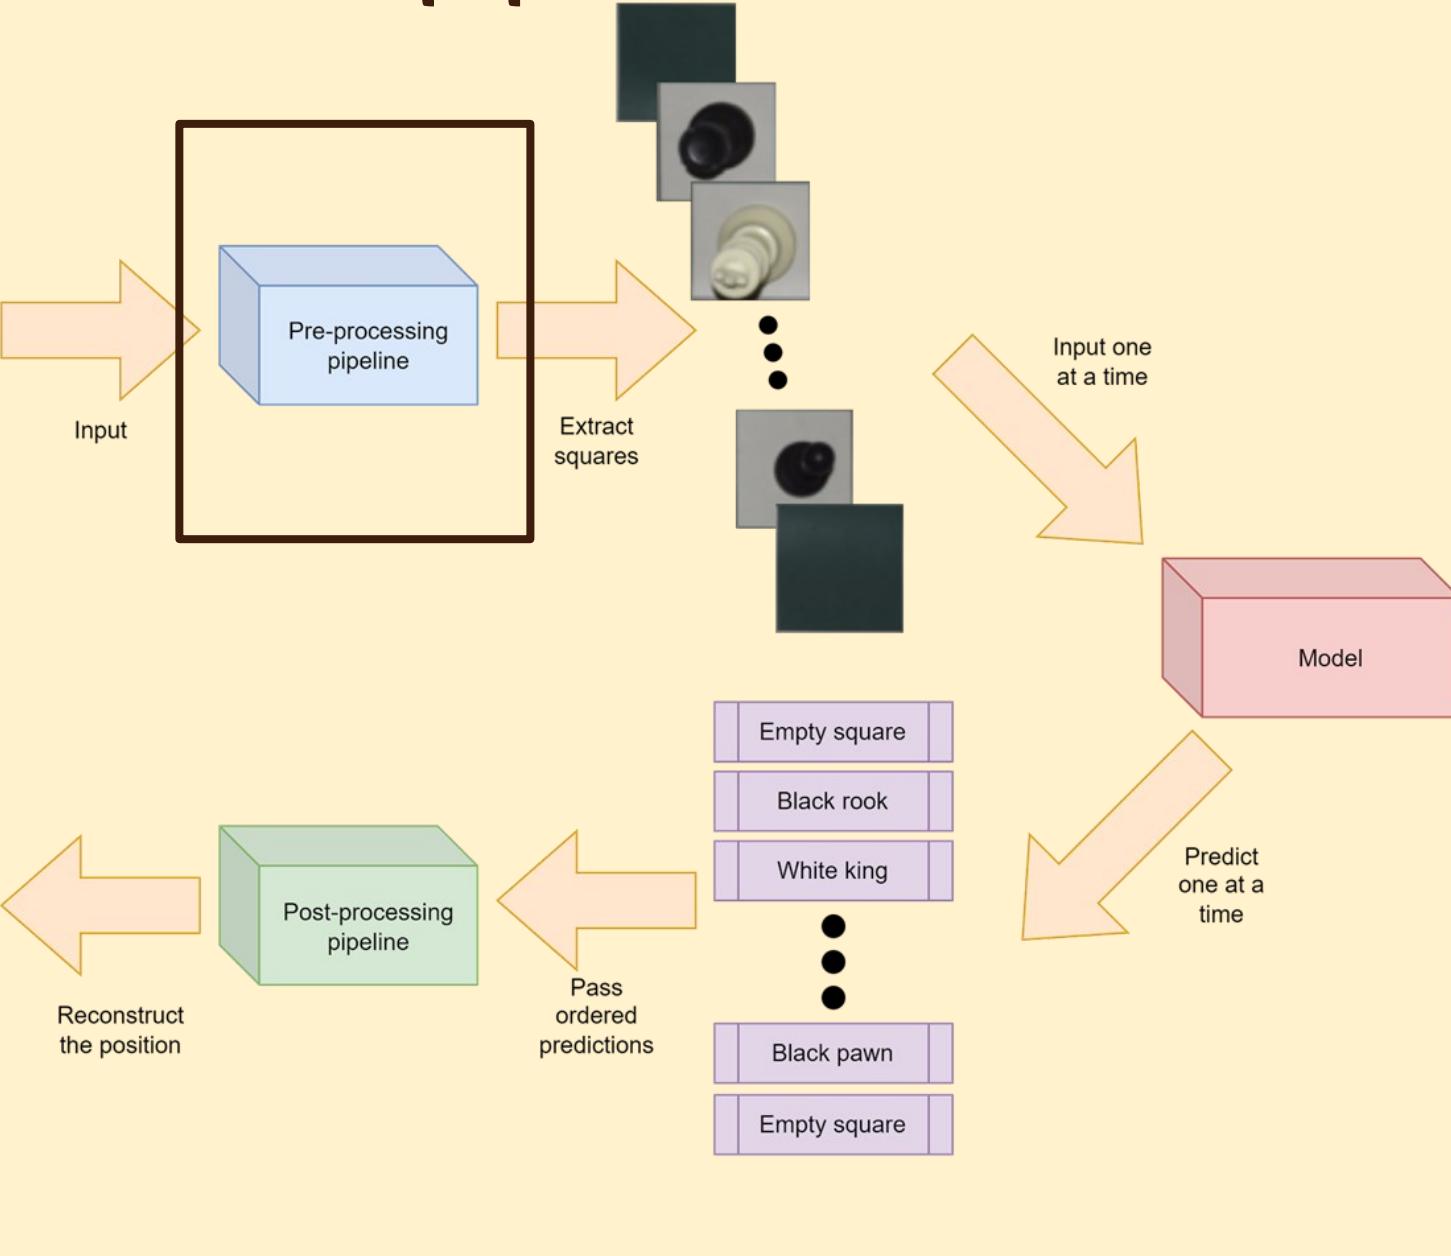
1- Generating full positions leads to more realistic lighting condition distributions, since a square must now be part of a bigger picture which is the full board.

2- All Generated positions are legal. In that sense, subsequent iterations of this work could incorporate constraints stemming from chess rules.

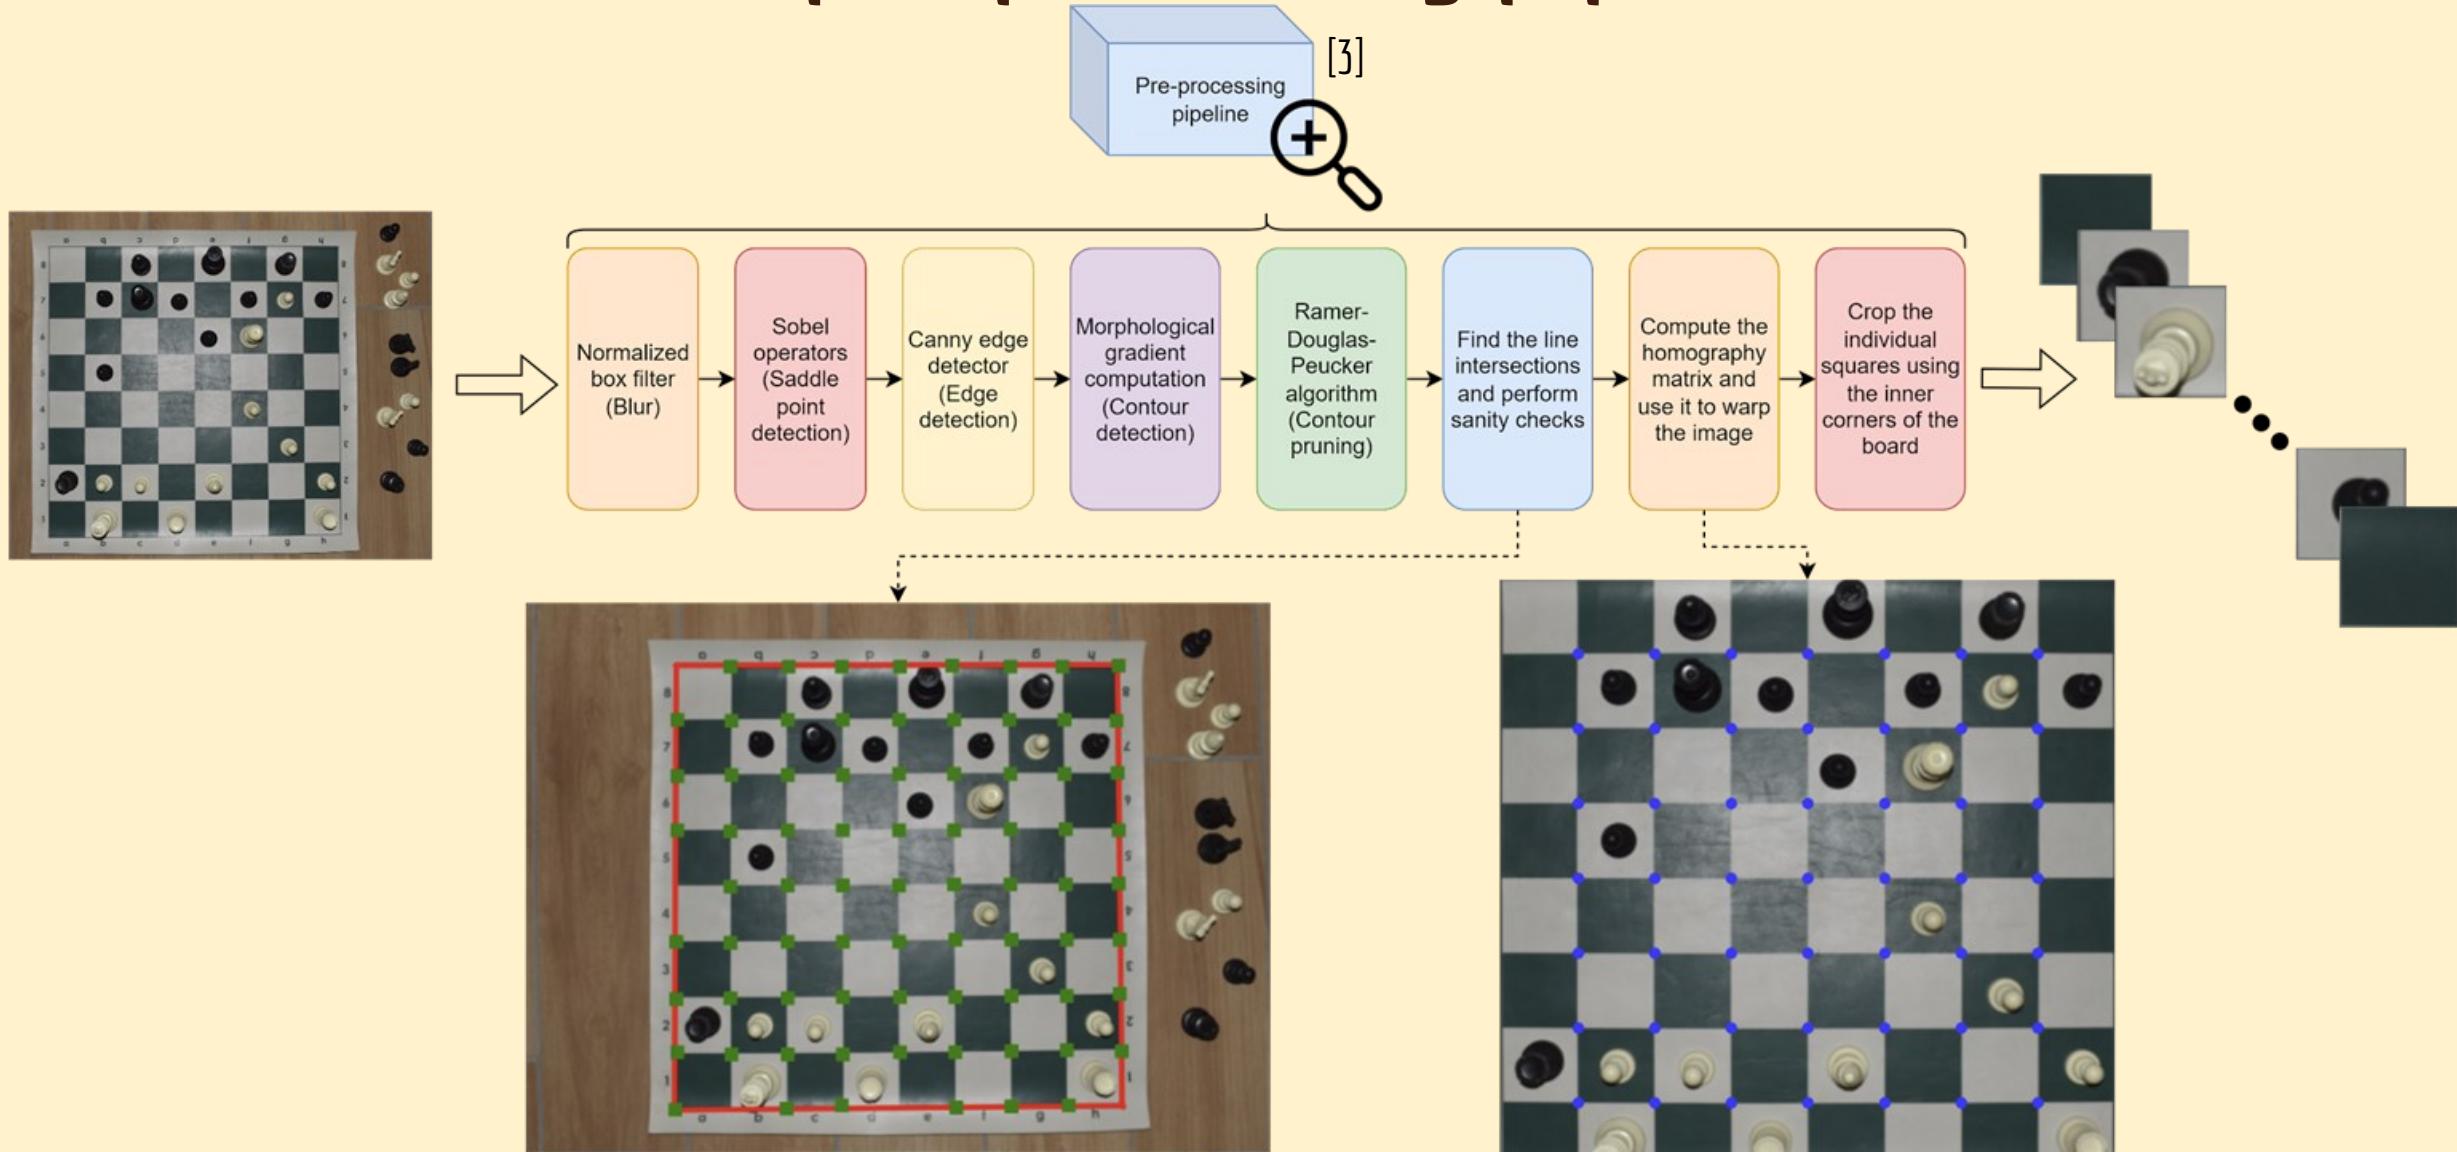
# End-to-end pipeline overview



# Full end-to-end pipeline overview



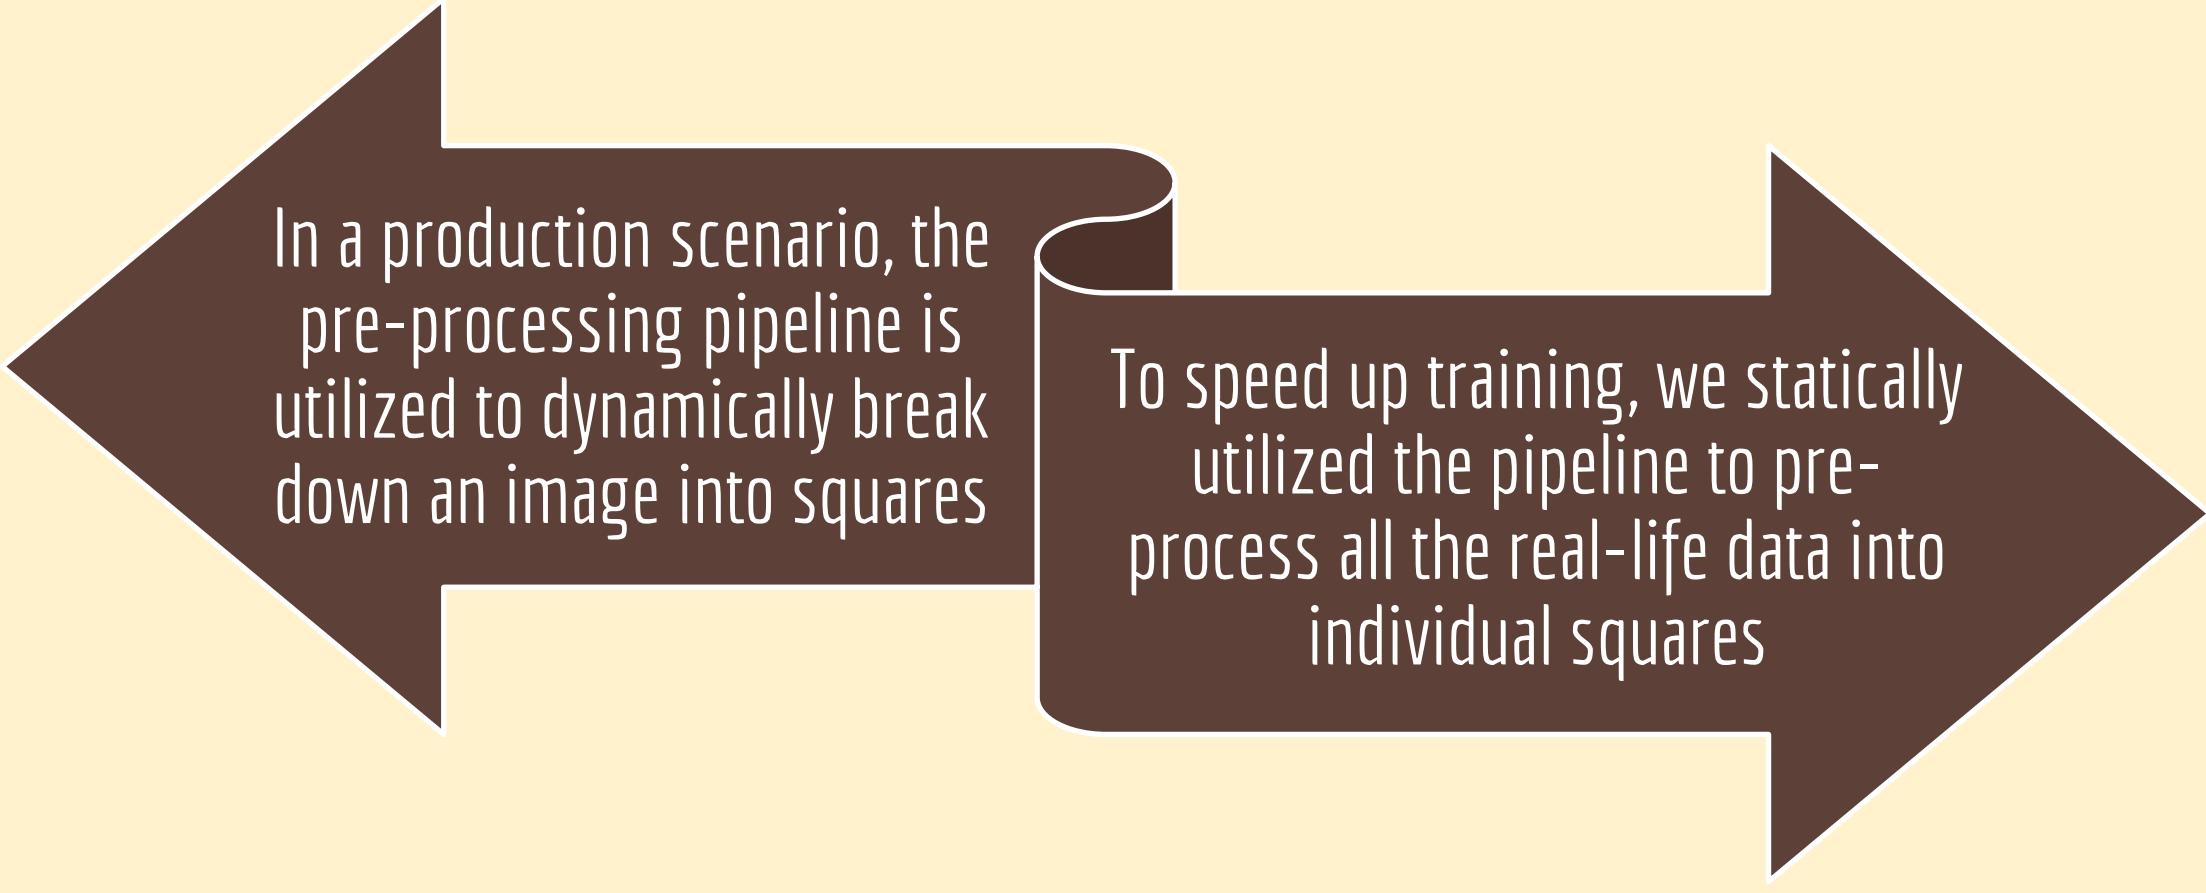
# The pre-processing pipeline



# Interlude: Dynamic VS Static pre-processing

Time required to pre-process 1 image: A few seconds

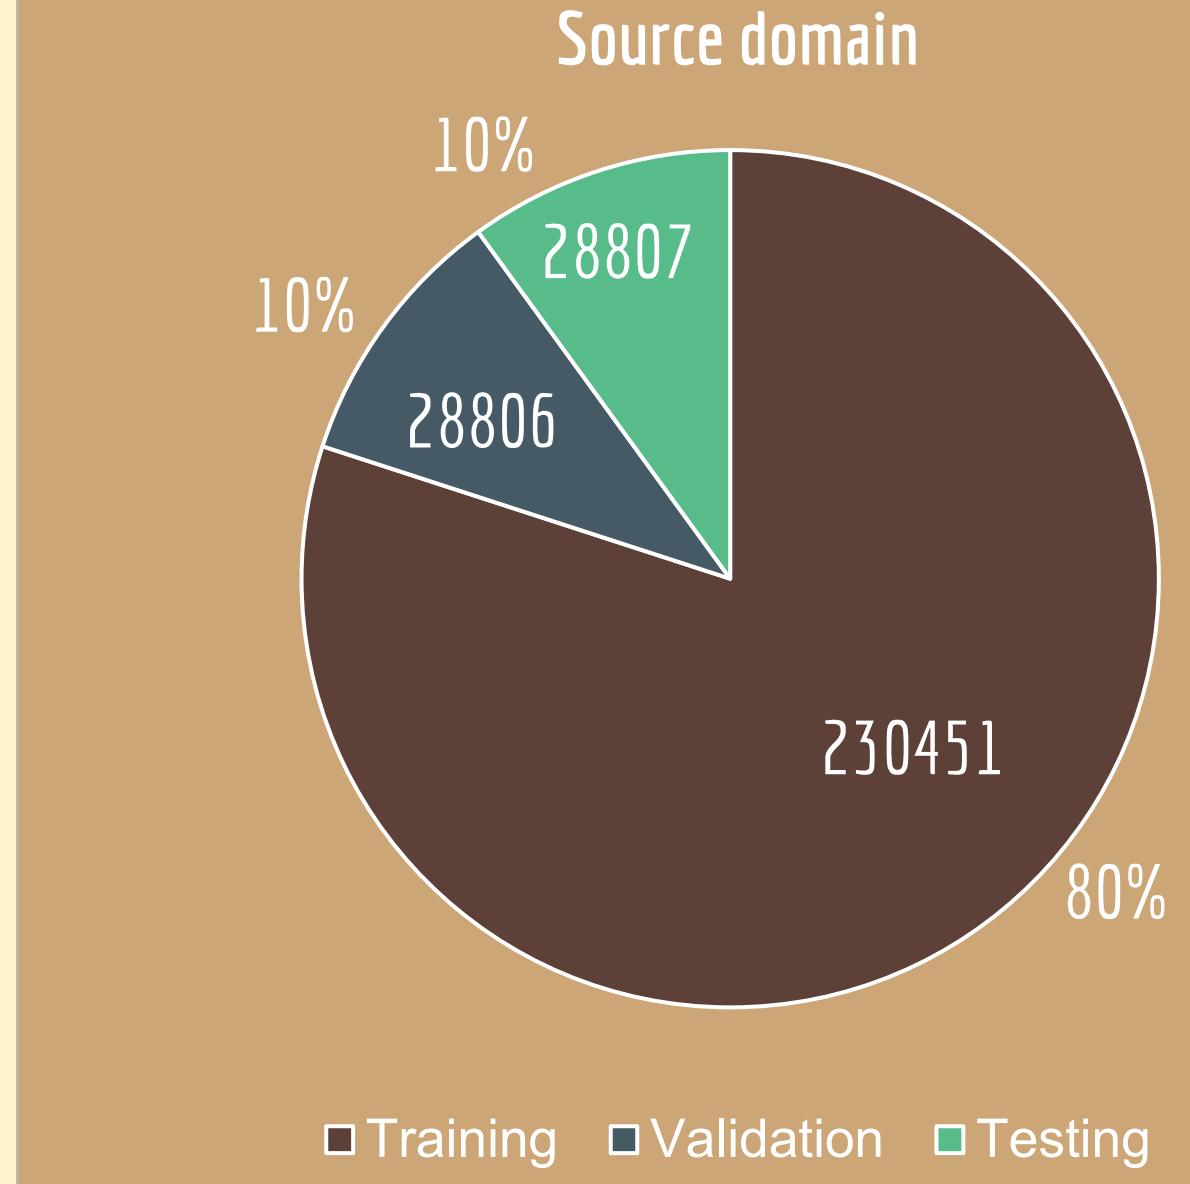
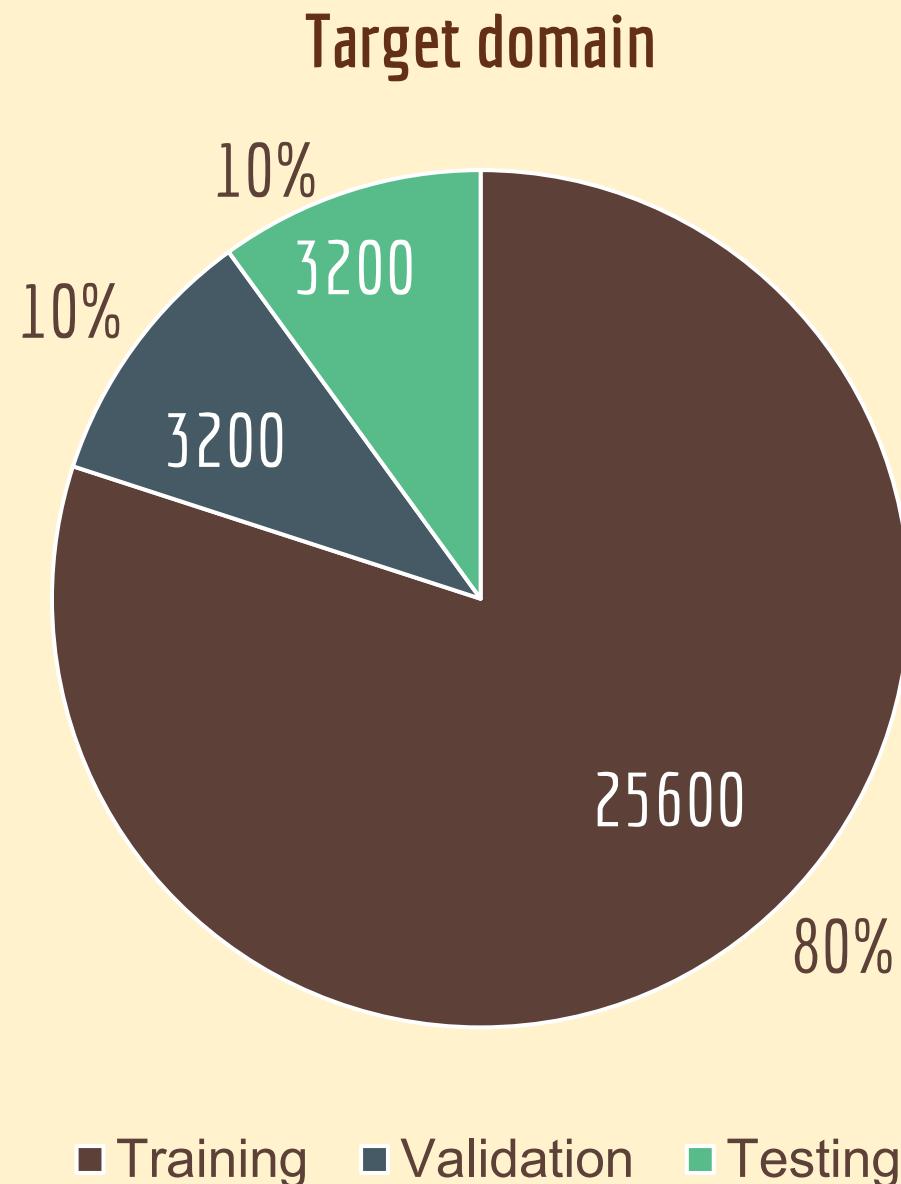
This can slow down the training drastically



In a production scenario, the pre-processing pipeline is utilized to dynamically break down an image into squares

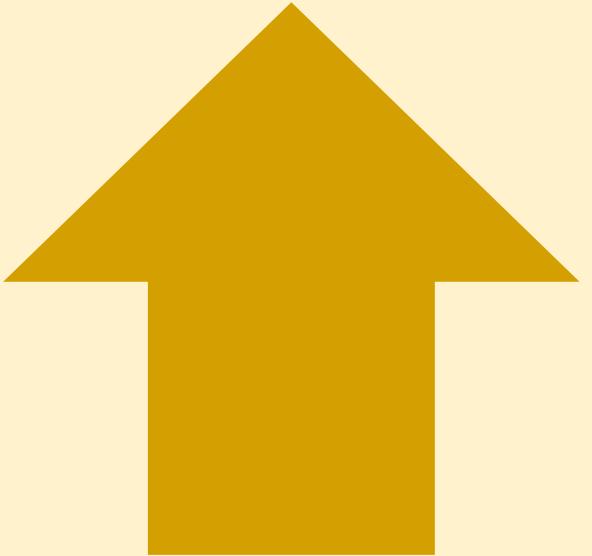
To speed up training, we statically utilized the pipeline to pre-process all the real-life data into individual squares

# Interlude: Dataset splits

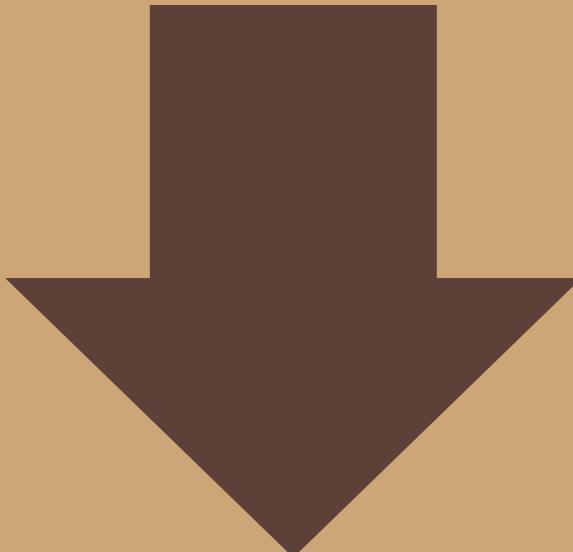




# Interlude: Note about target domain labelling



The training set of the target domain was completely unlabelled throughout the experiment



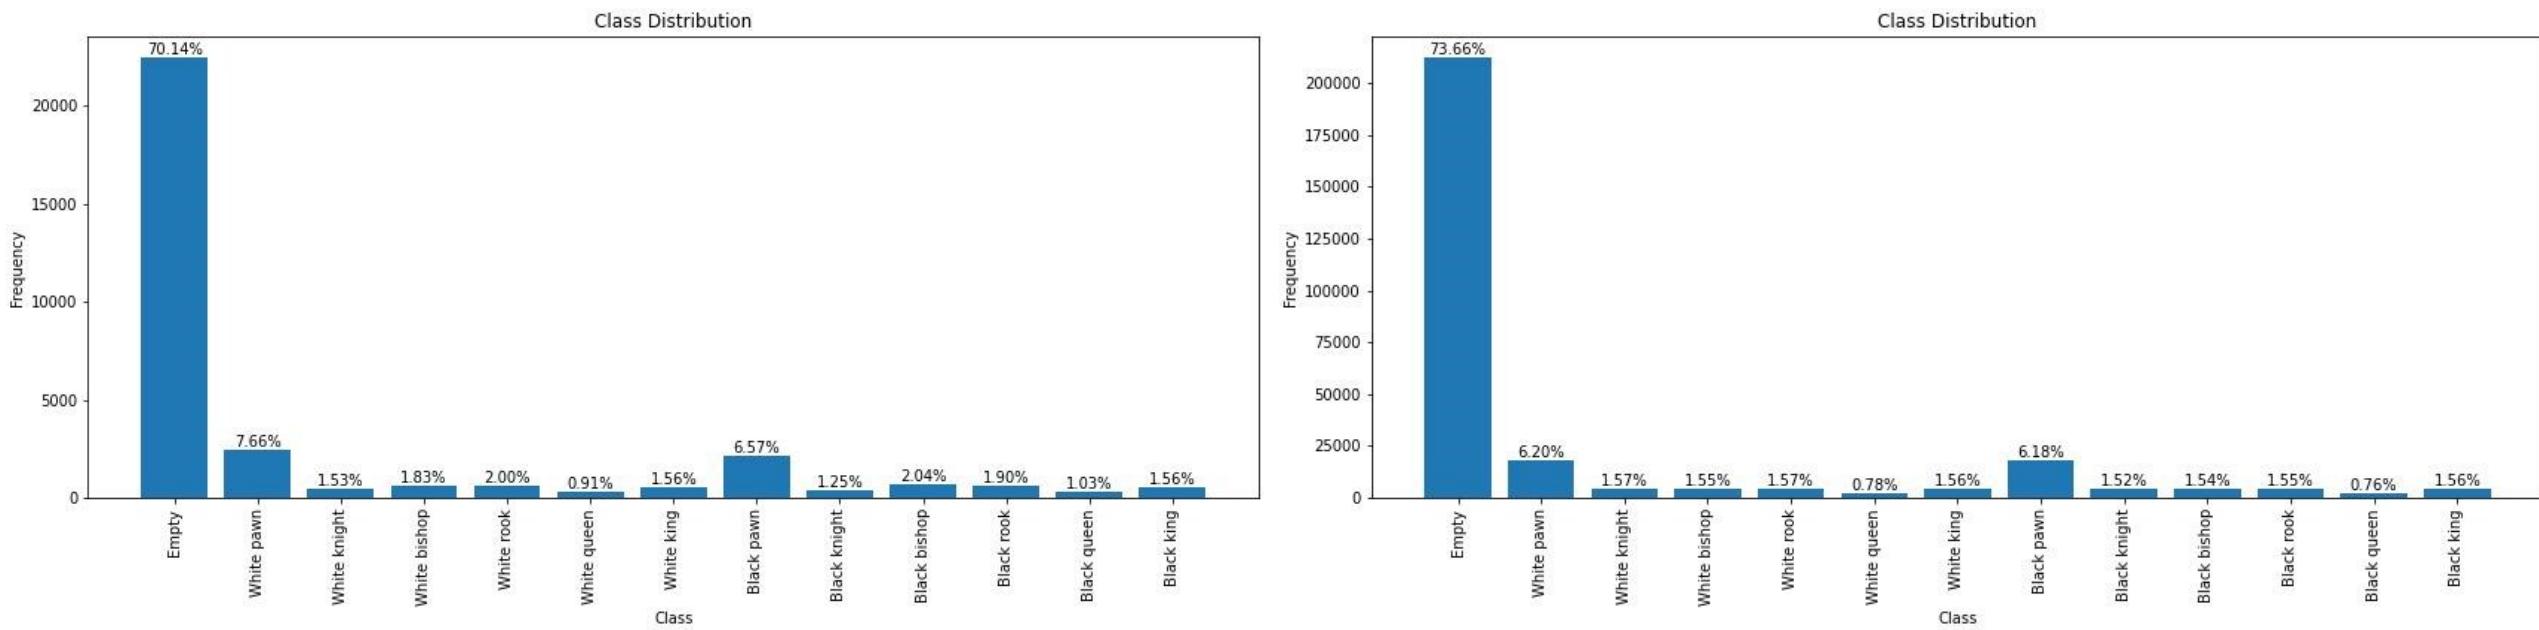
We used the validation set's labels to perform hyperparameter tuning, and for the testing set to plot metrics such as weighted F1 score & confusion matrices



# Interlude: Dataset distributions

Target domain

Source domain

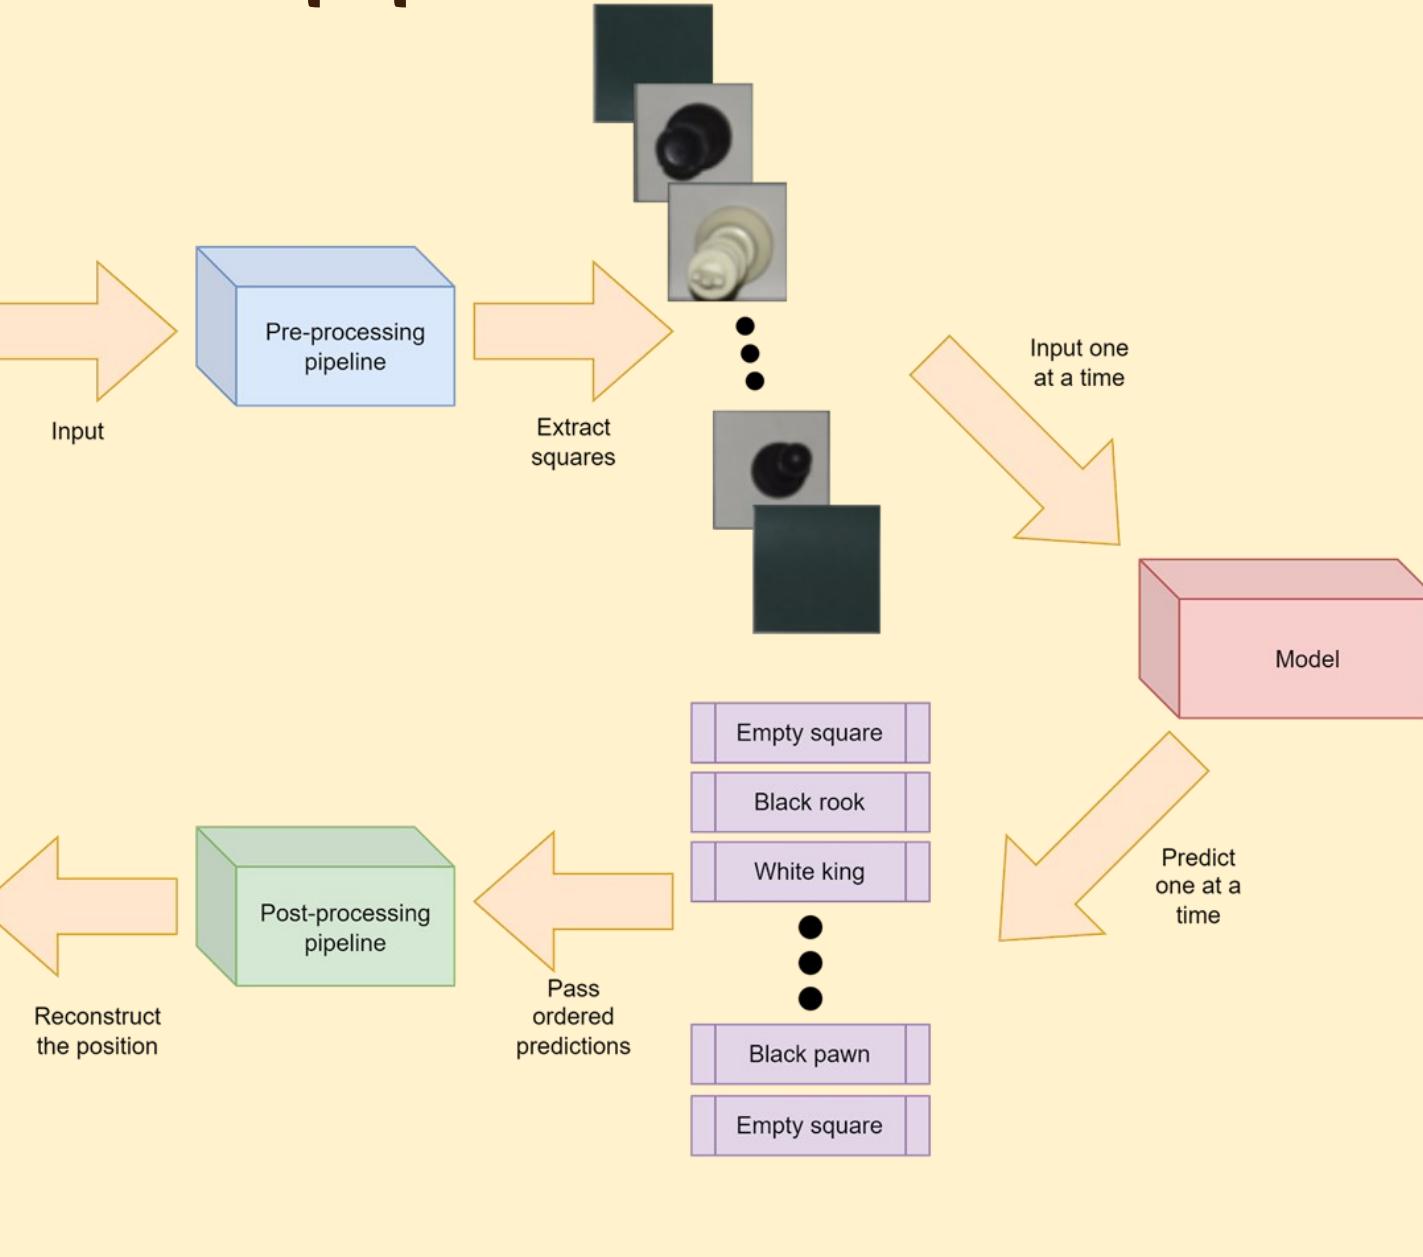


Empty squares dominate

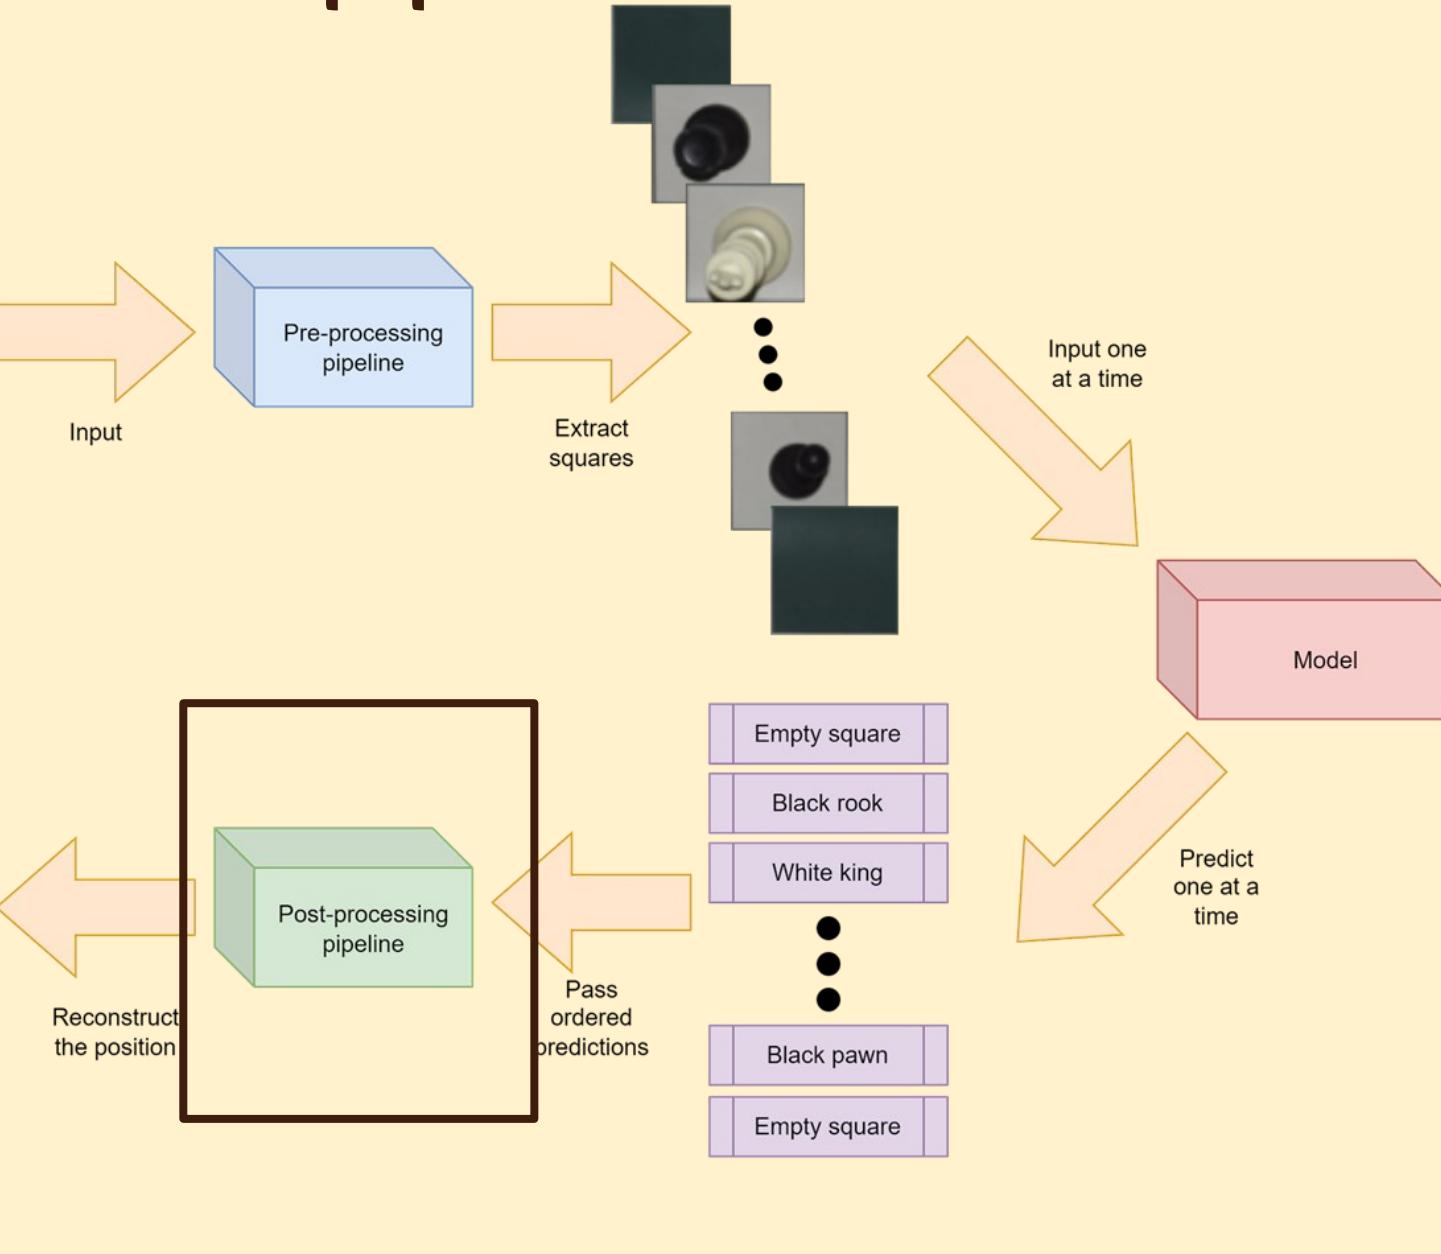
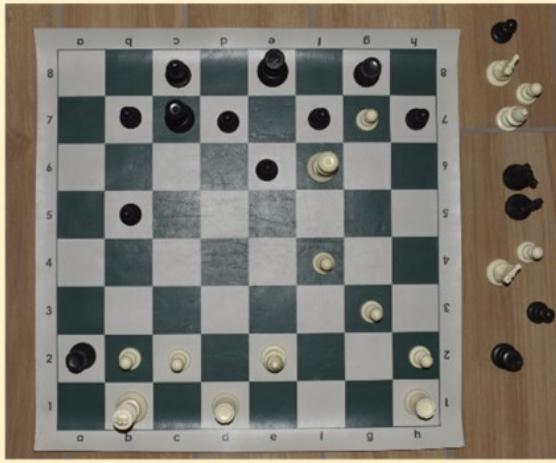
Utilized PyTorch's `WeightedRandomSampler` with replacement to re-balance the minibatch proportions

Augmented the data (Random flips and brightness jitters) to avoid just repeating minority class samples

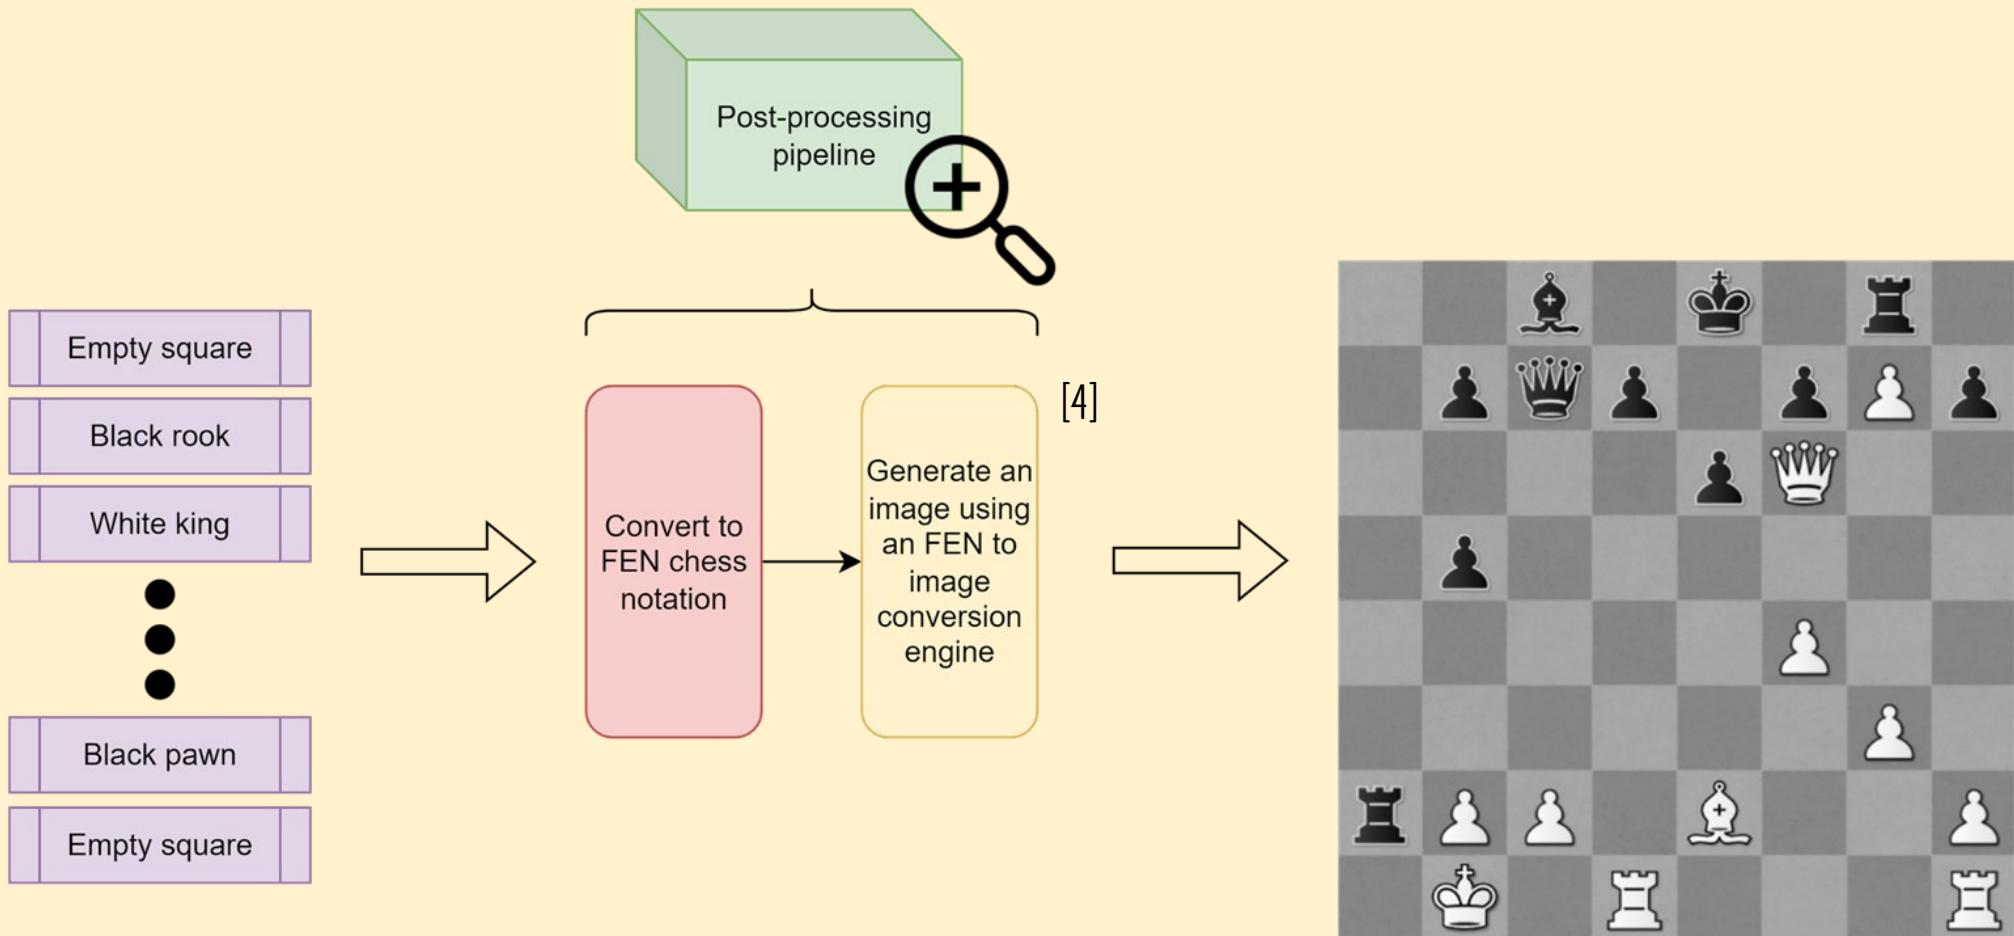
# End-to-end pipeline overview



# End-to-end pipeline overview

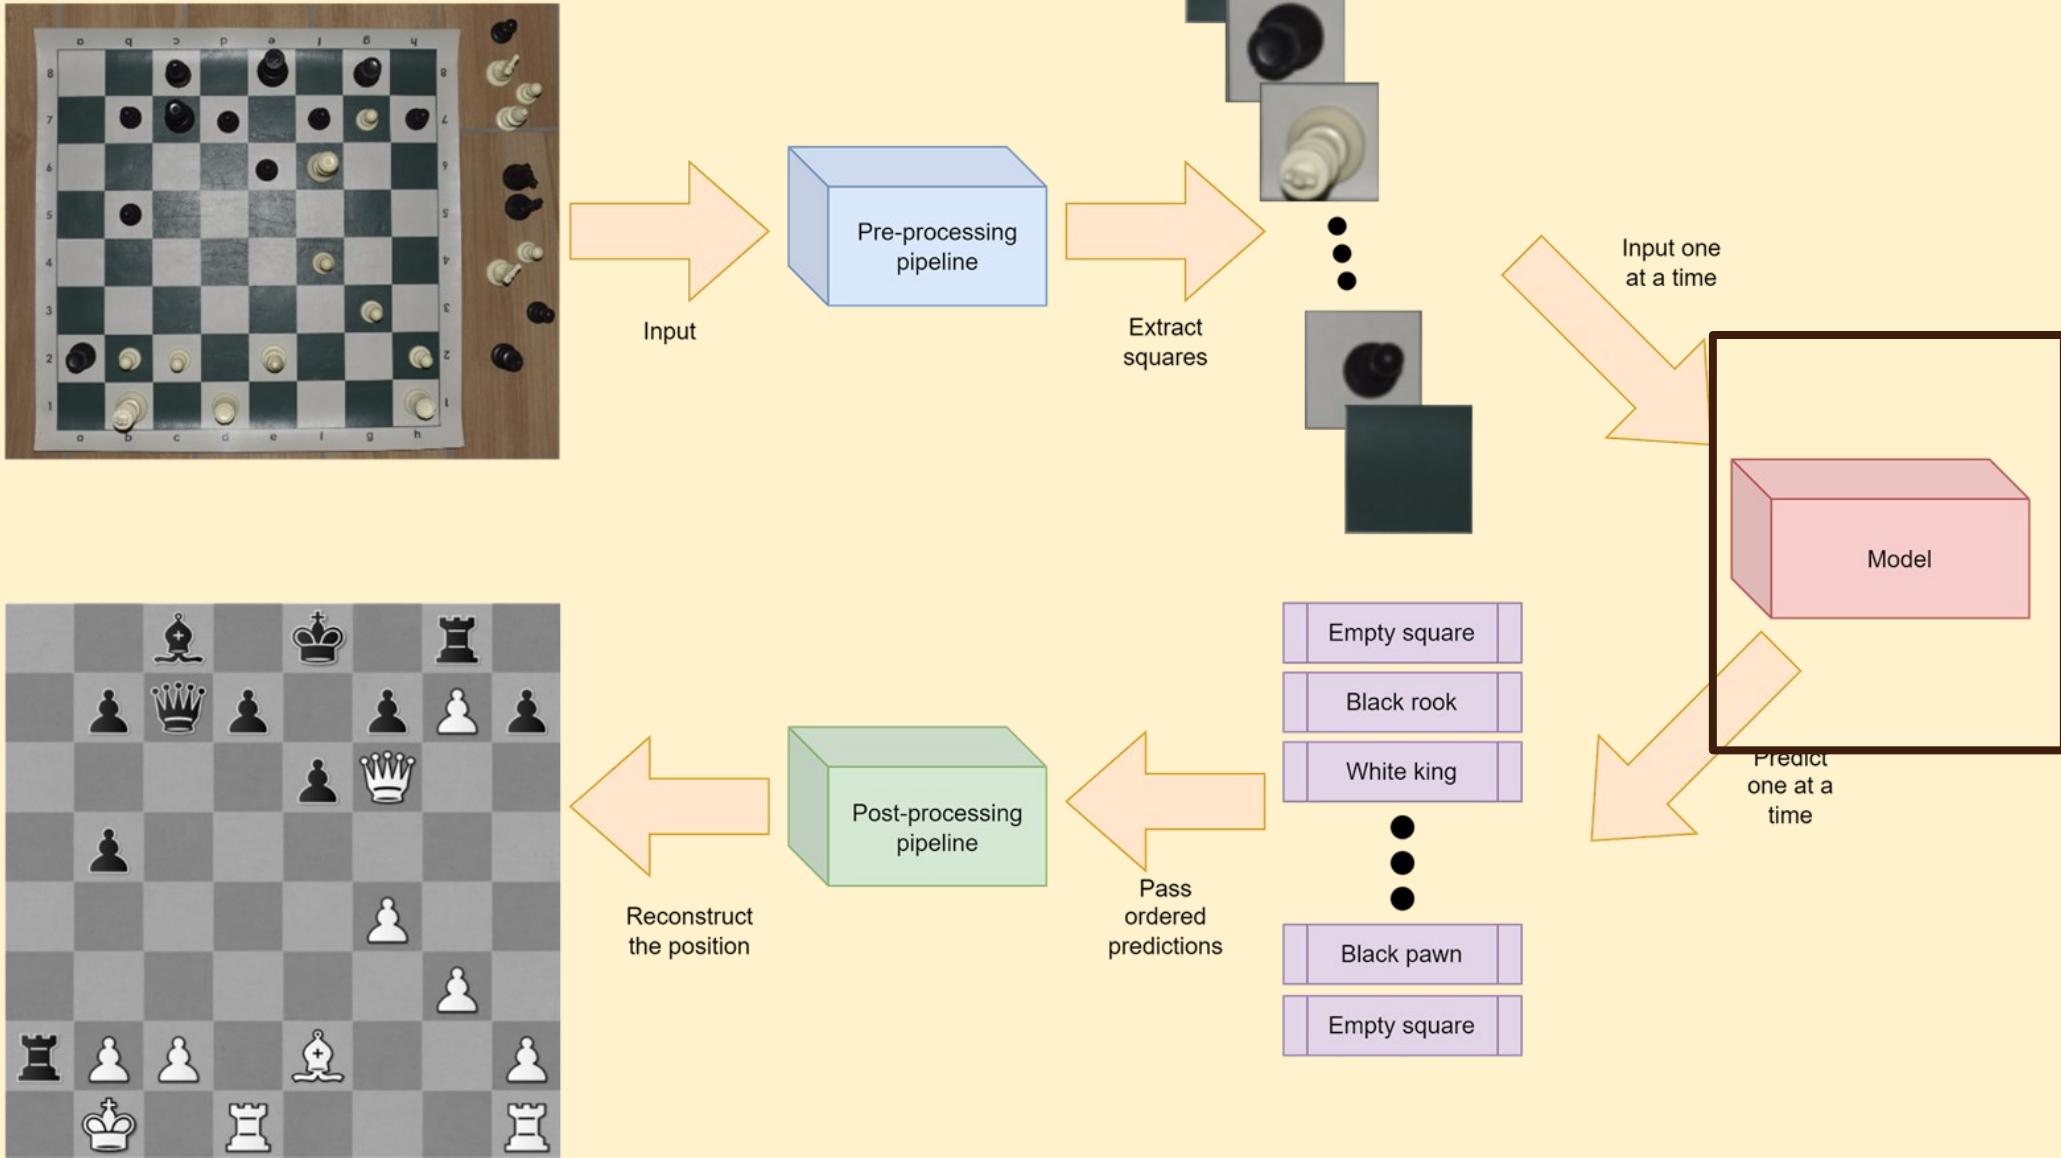


# The post-processing pipeline



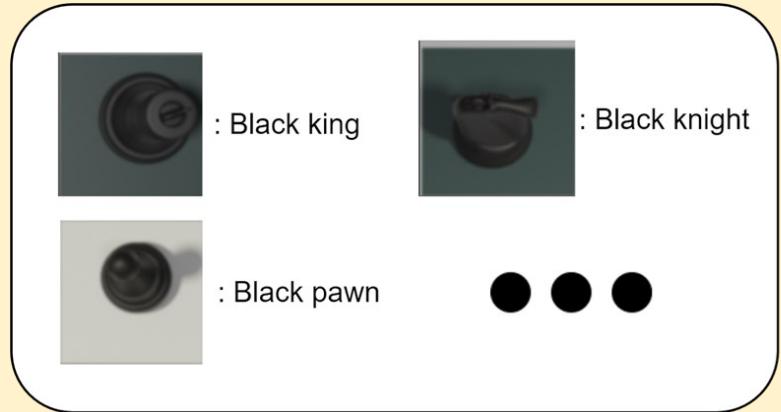
[4] “Fen to image,” Chessvisionai RSS, <https://chessvision.ai/docs/tools/fen2image/> (accessed Jun. 1, 2023).

# End-to-end pipeline overview

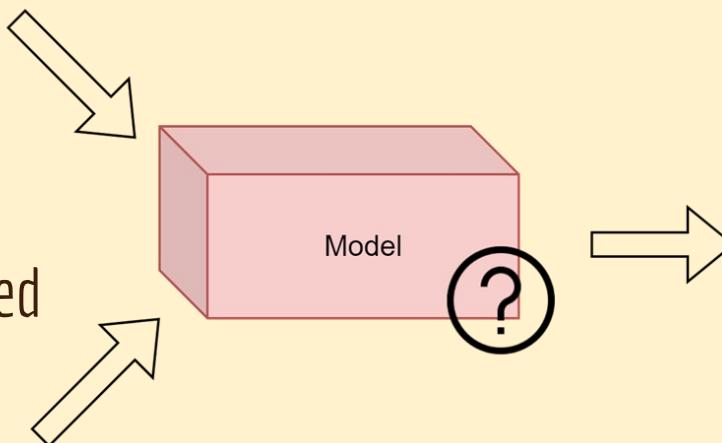
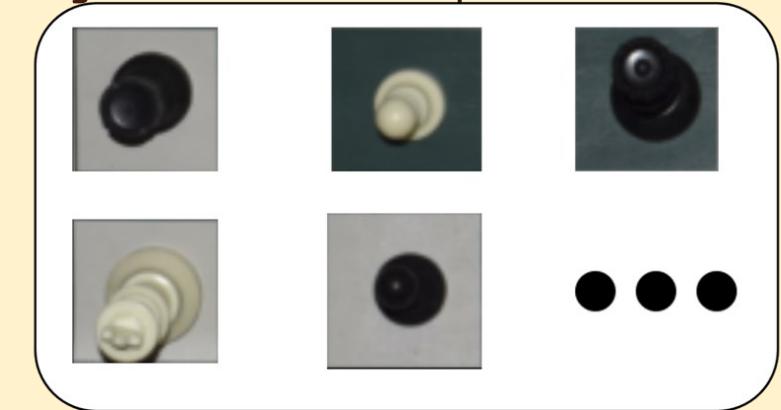


# The model: Architectures

Source domain: Generated data, labelled



Target domain: Real life data, unlabelled

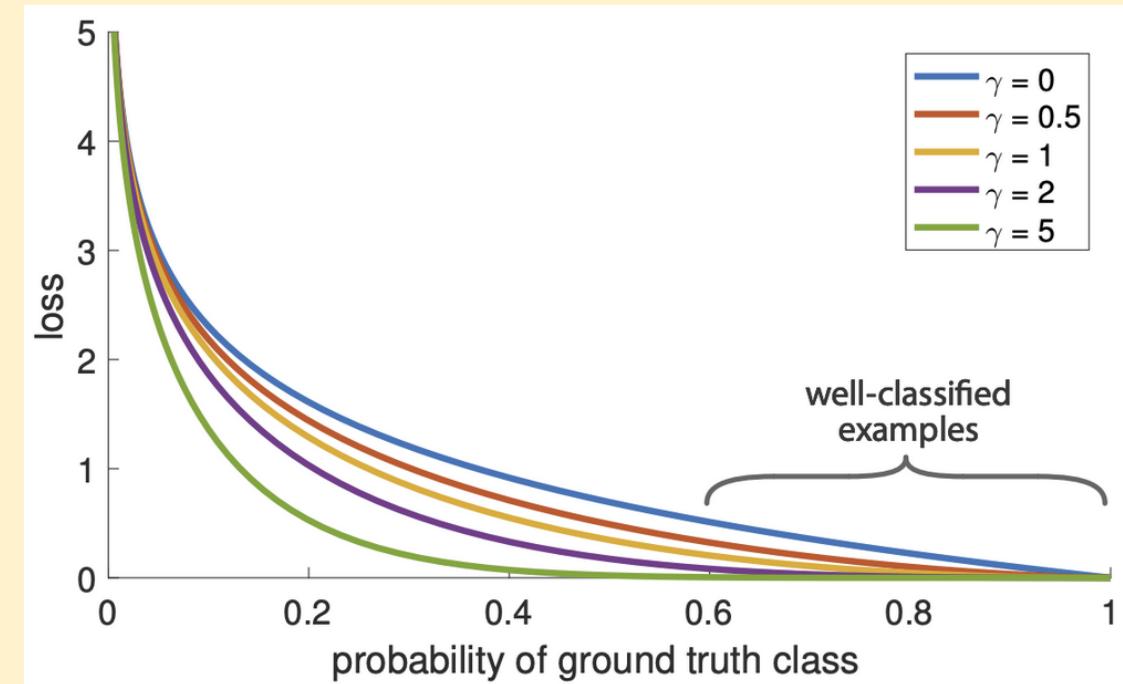


What domain adaptation  
technique to use in order to get  
good accuracies on the target  
domain?

# Interlude: Loss functions

Cross Entropy =  $-\log(p_i)$

Focal Loss<sup>[5]</sup> =  $-(1 - p_i)^\gamma \log(p_i)$



CORAL Loss<sup>[6]</sup> =  $\left\| C_s - C_t \right\|_F^2$  with  $C(X) = \frac{1}{n-1} (X - \bar{X})^T (X - \bar{X})$

[5] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, 'Focal Loss for Dense Object Detection', *arXiv [cs.CV]*. 2018.

[6] B. Sun, J. Feng, and K. Saenko, 'Correlation Alignment for Unsupervised Domain Adaptation', *CoRR*, vol. abs/1612.01939, 2016.

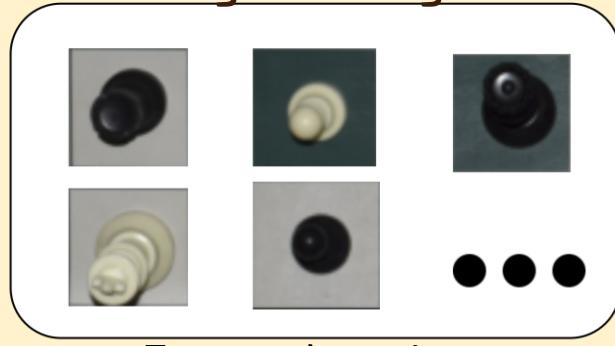
# Architecture 1: VGG16 without DA (BASE-source)

During training

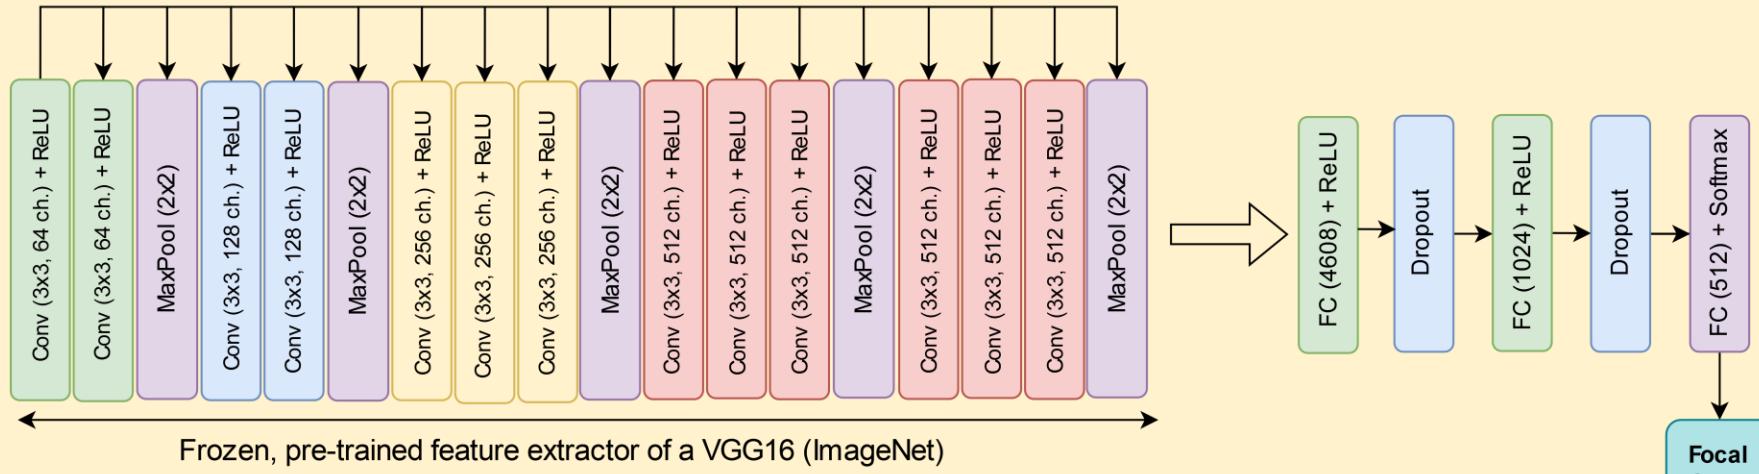


Source domain

During testing



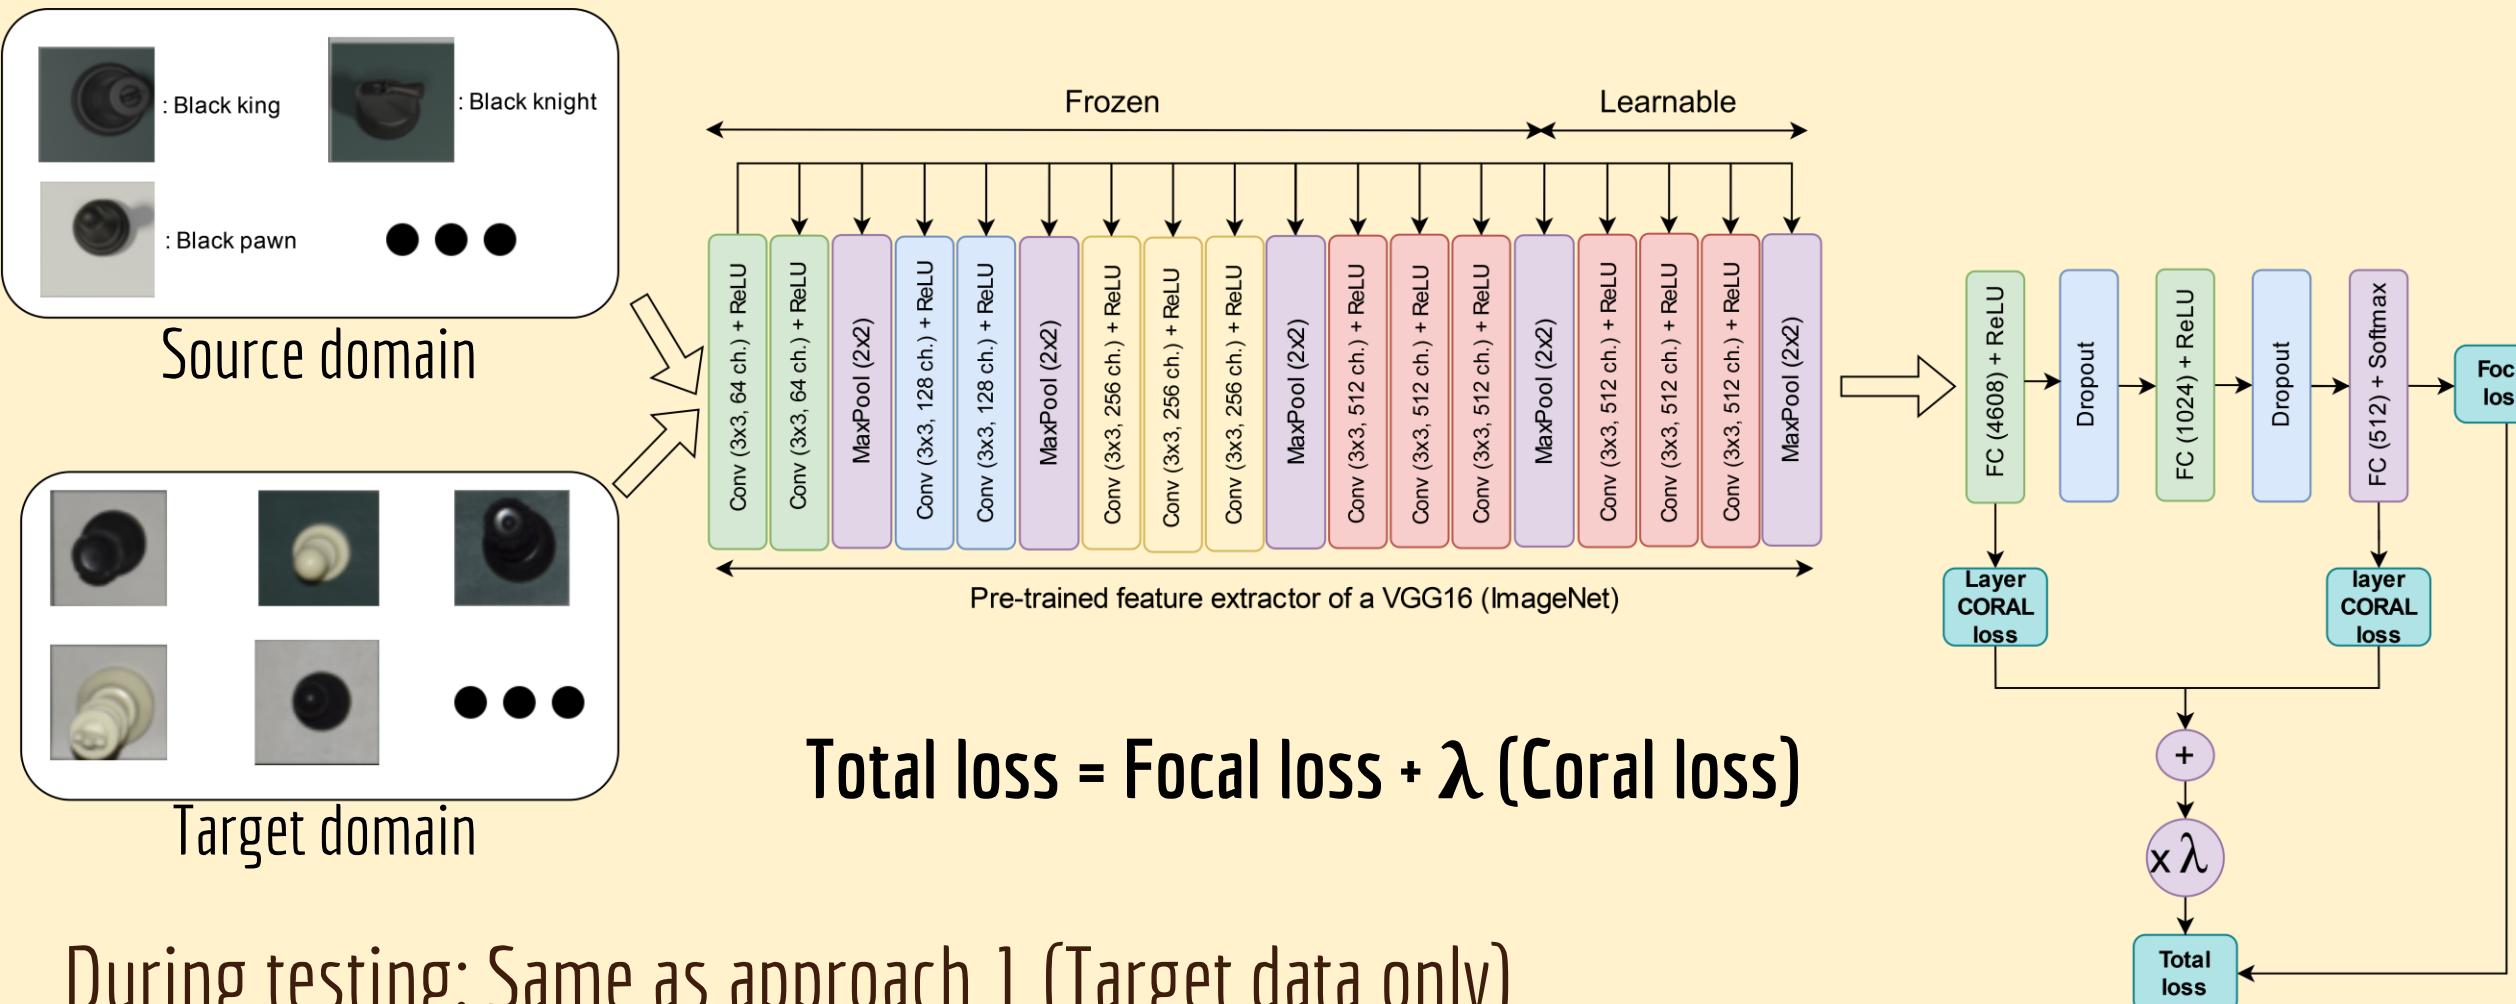
Target domain



Frozen, pre-trained feature extractor of a VGG16 (ImageNet)

# Architecture 2: Correlation alignment (CORAL)

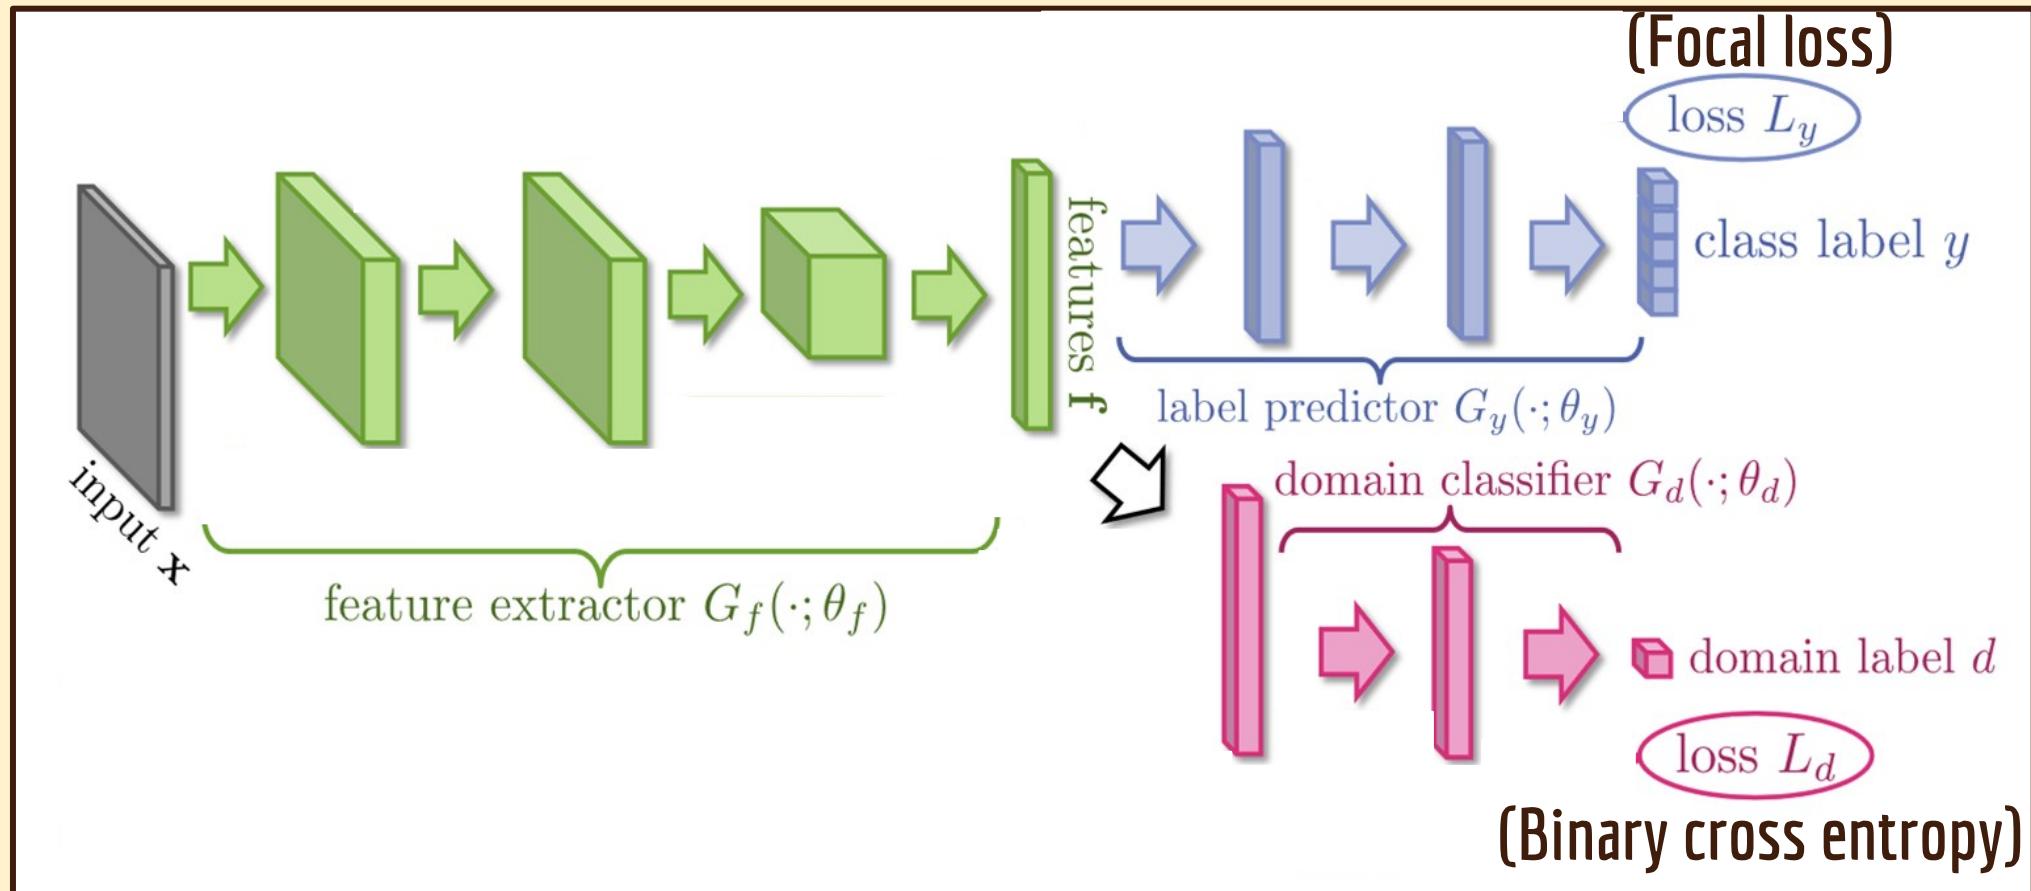
During training



During testing: Same as approach 1 (Target data only)

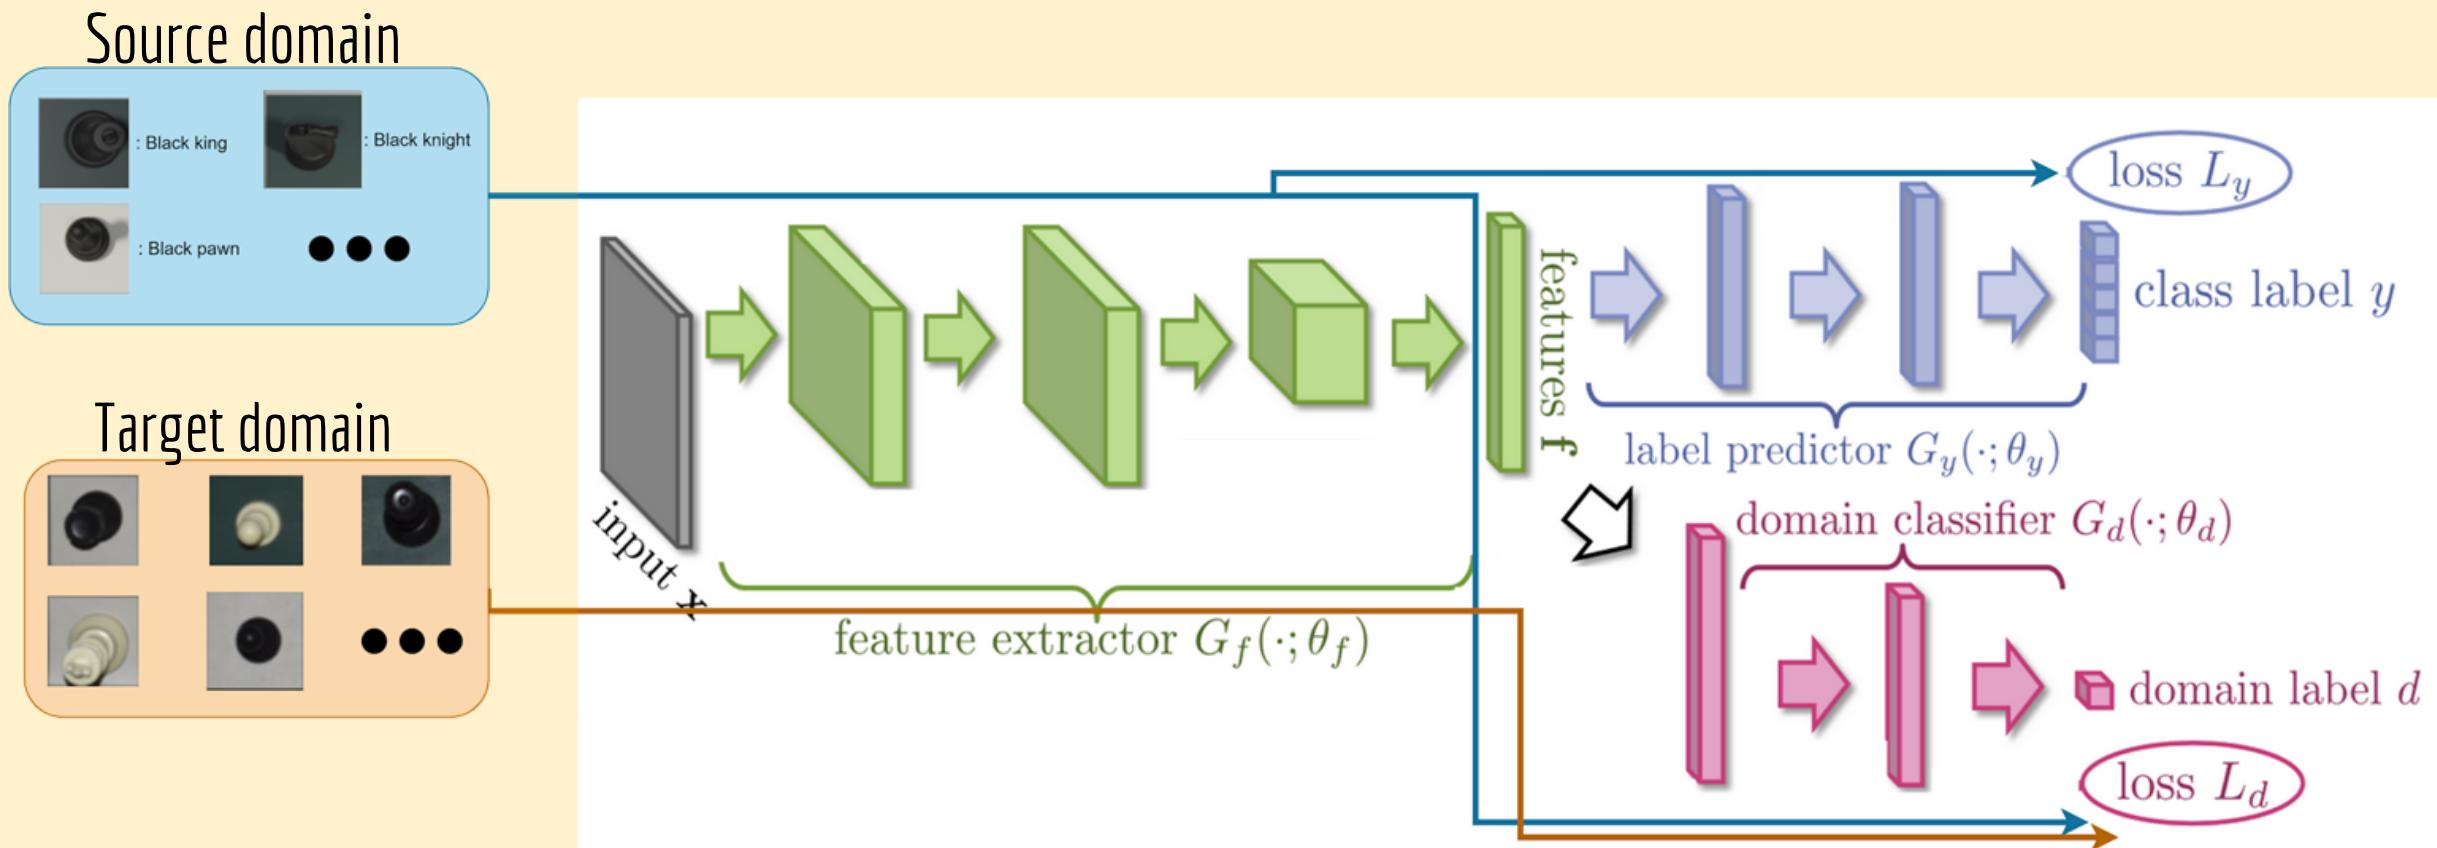
# Architecture 3: Adversarial approach (DANN)

During training



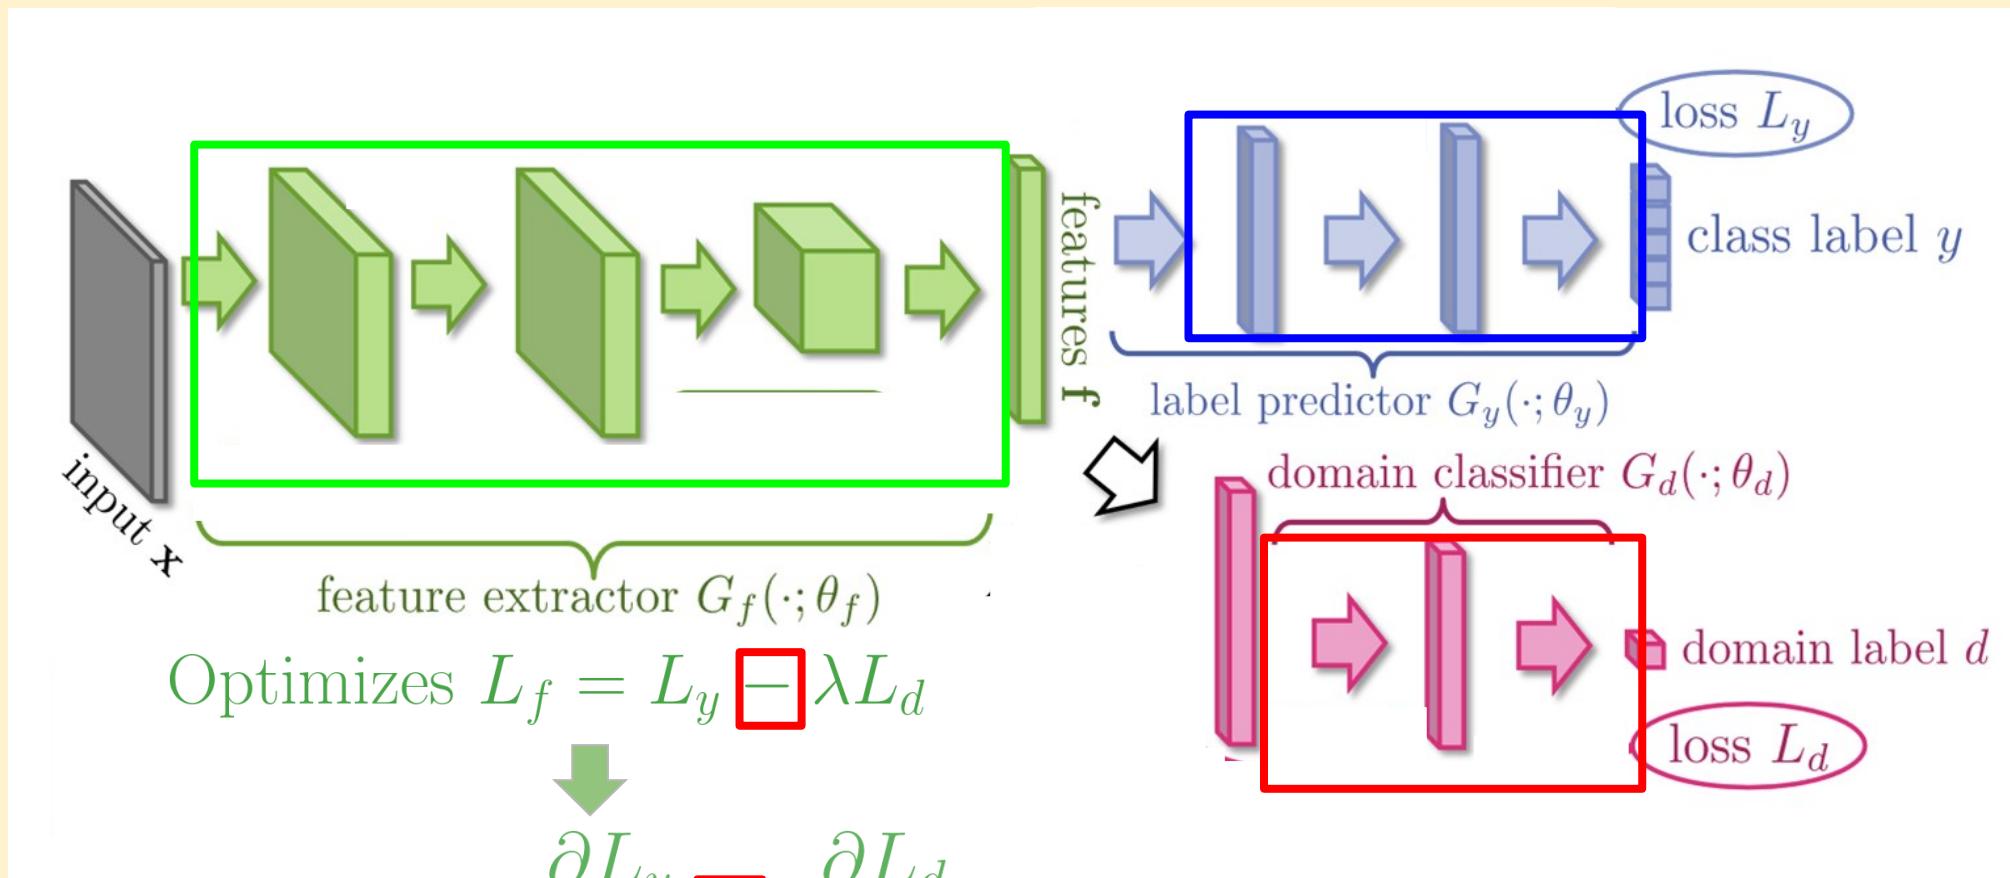
# Architecture 3: Adversarial approach (DANN)

During training: Forward pass



# Architecture 3: Adversarial approach (DANN)

During training: Backward pass (1)



$$\theta_f = \theta_f - \frac{\partial L_y}{\partial \theta_f} \square \lambda \frac{\partial L_d}{\partial \theta_f}$$

Optimizes  $L_y$

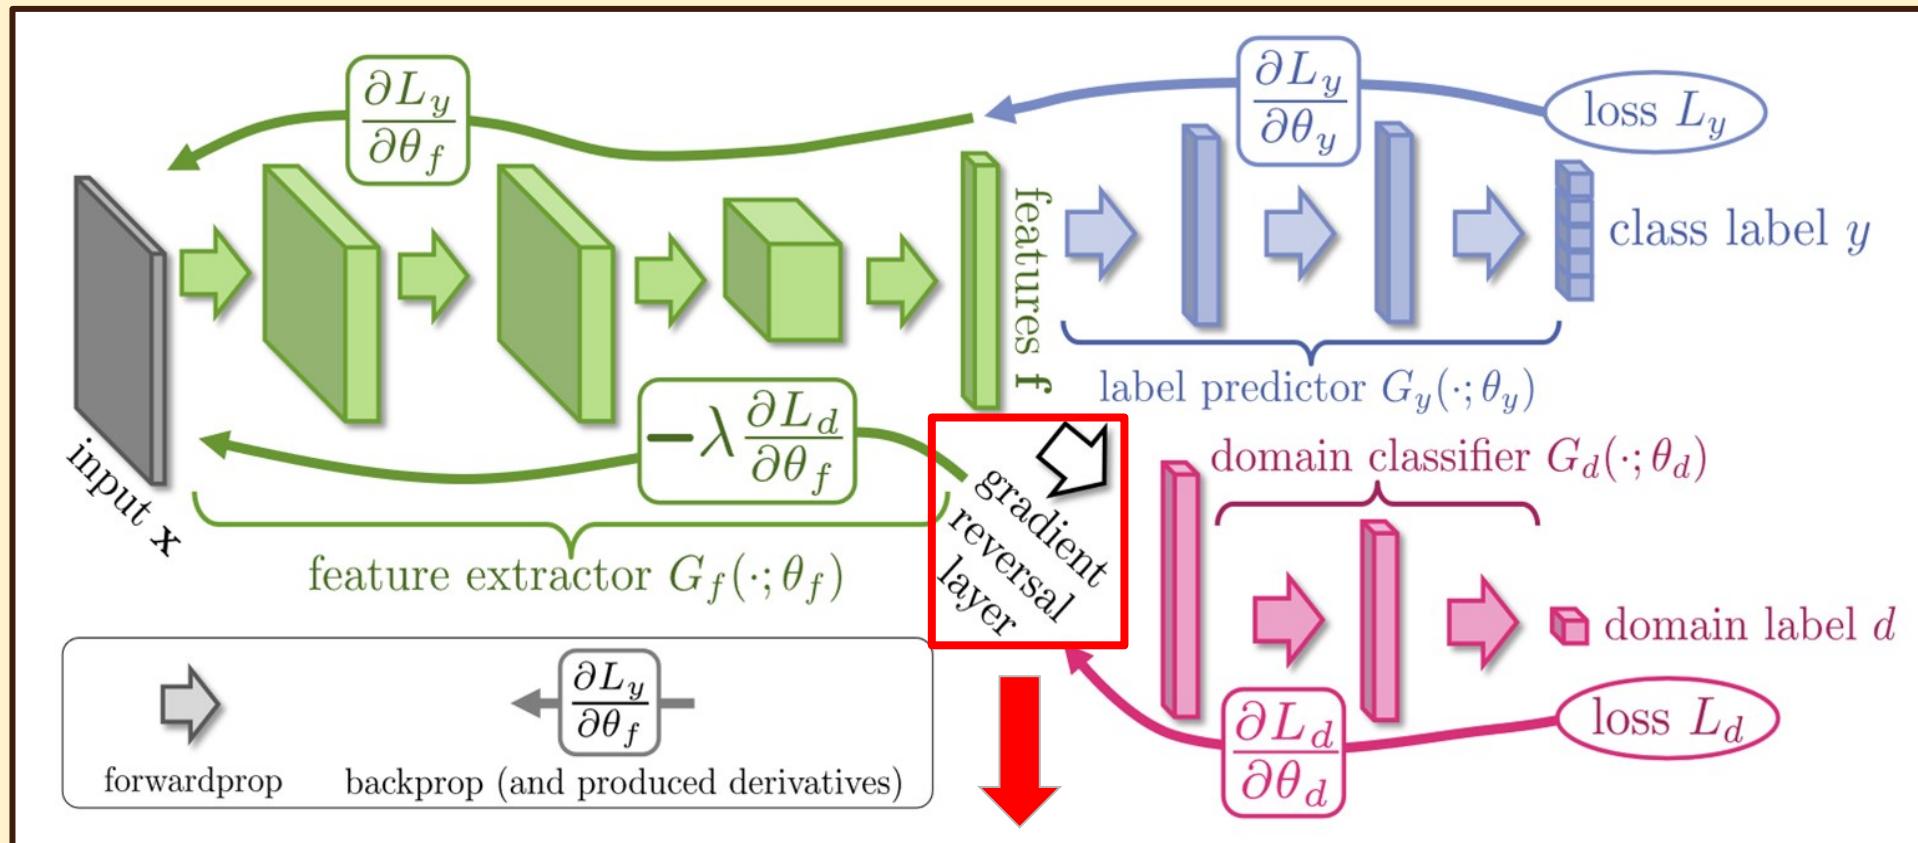
$$\theta_y = \theta_y - \frac{\partial L_y}{\partial \theta_y}$$

Optimizes  $L_d$

$$\theta_d = \theta_d - \frac{\partial L_d}{\partial \theta_d}$$

# Architecture 3: Adversarial approach (DANN)

During training: Backward pass (2)



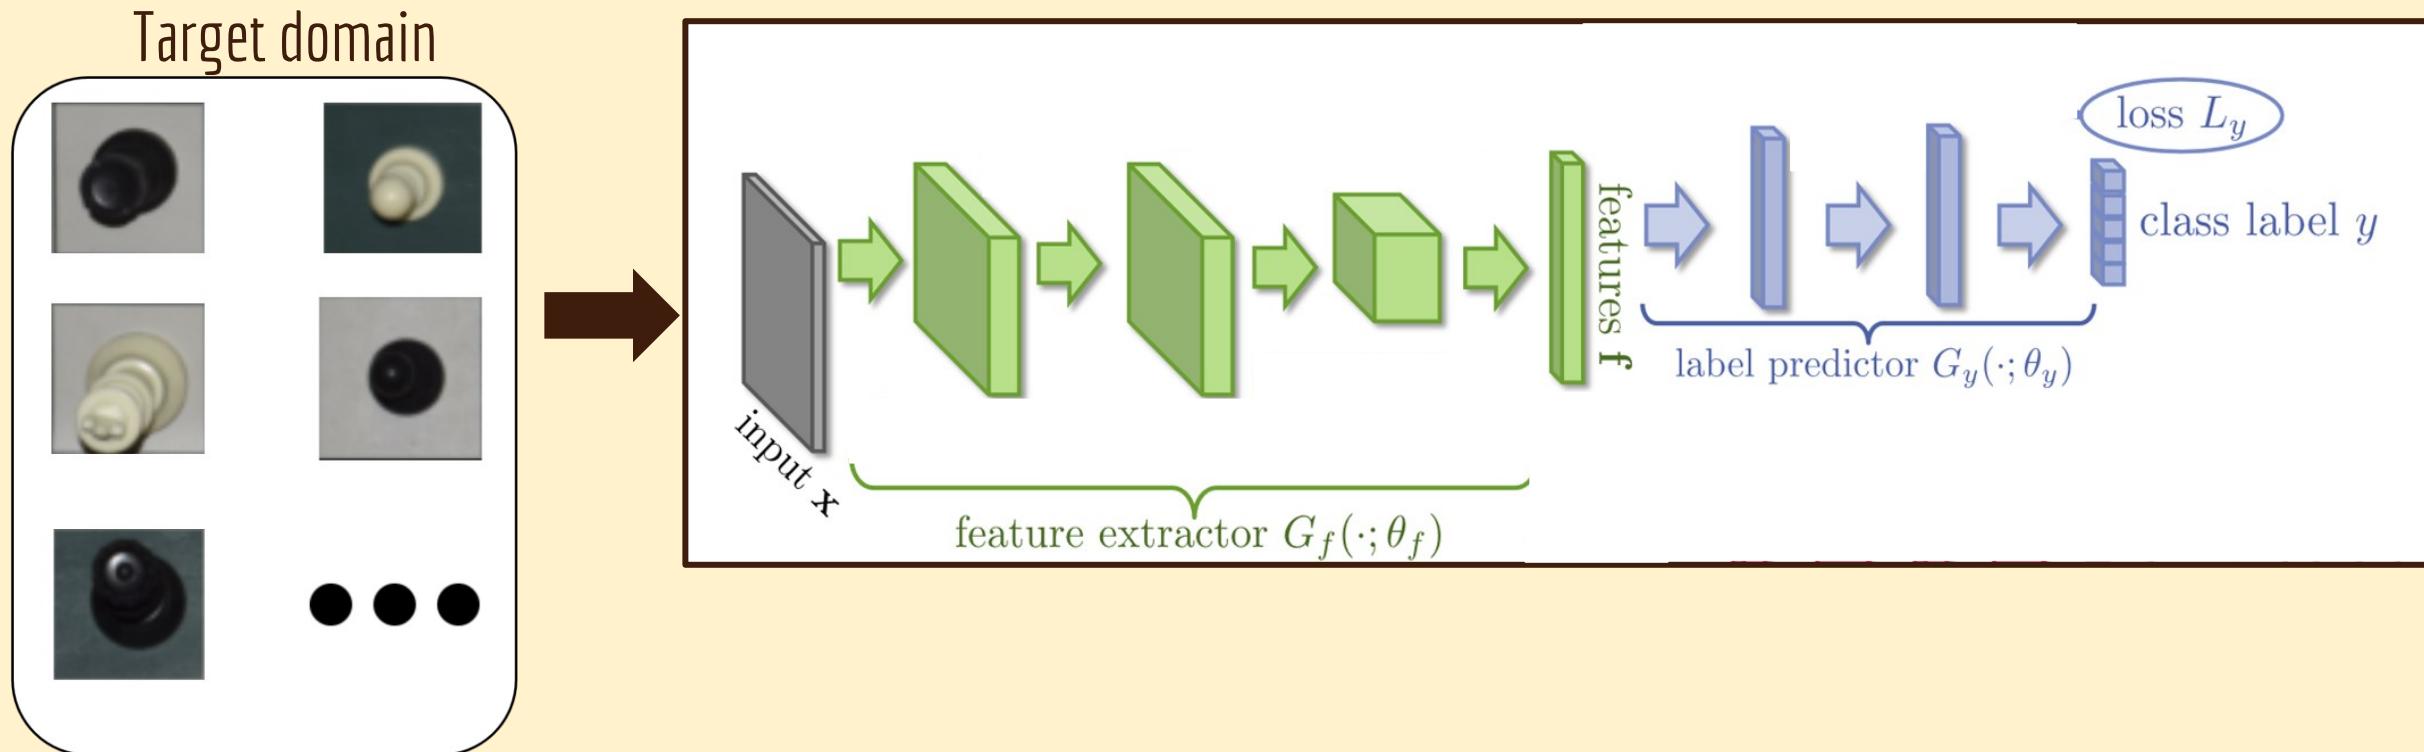
Gradient reversal layer:

- Identity function during forward pass
- Multiplies by  $-\lambda$  during the backward pass

$$\theta_f = \theta_f - \frac{\partial L_y}{\partial \theta_f} + \lambda \frac{\partial L_d}{\partial \theta_f}$$

# Architecture 3: Adversarial approach (DANN)

During testing



Note: The domain discriminator can be discarded during testing

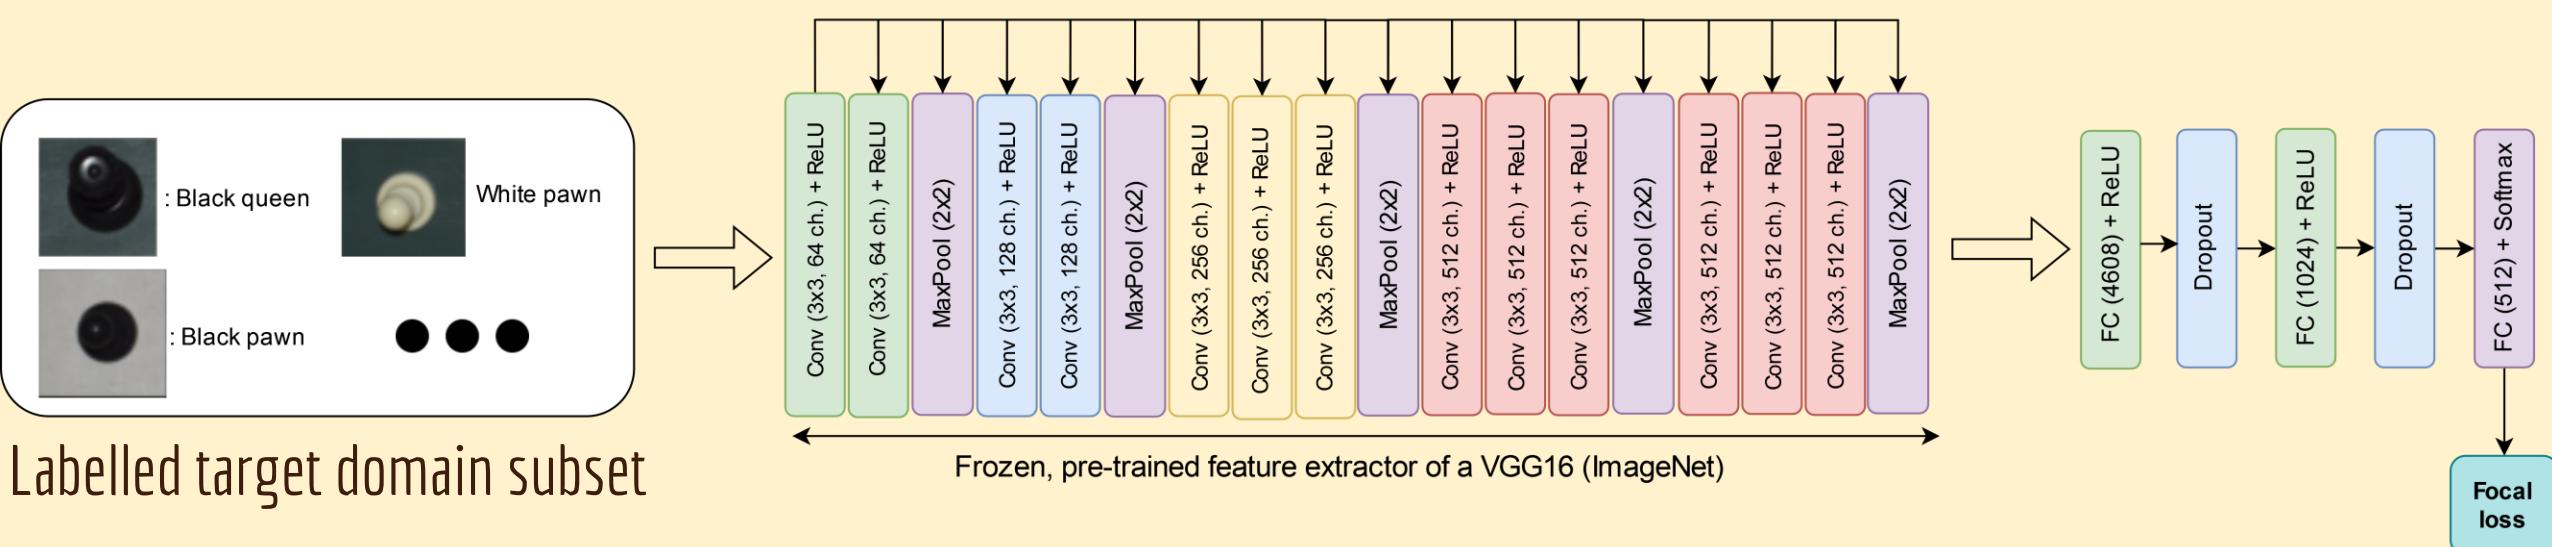
# Architecture 4: VGG16 trained on a labelled target domain subset of the validation data (BASE-target)

We used a labelled validation subset of 3200 target domain square images

But what if those were enough to train the model without DA?

Here, we split the full validation set into 80% training and 20% validation

We do not use the source domain dataset



# Hyperparameter tuning results - Base Models

Gamma	Dropout Rate	Best Validation Accuracy
2	0	59.7%
2	0.2	61.4%
2	0.5	56.7%
5	0	56.6%
5	0.2	56.9%
5	0.5	58.4%

Base-source

Gamma	Dropout Rate	Best Validation Accuracy
2	0	94.7%
2	0.2	93.3%
2	0.5	93.8%
5	0	94.7%
5	0.2	94.9%
5	0.5	93.8%

Base-target

# Hyperparameter tuning results - DA models

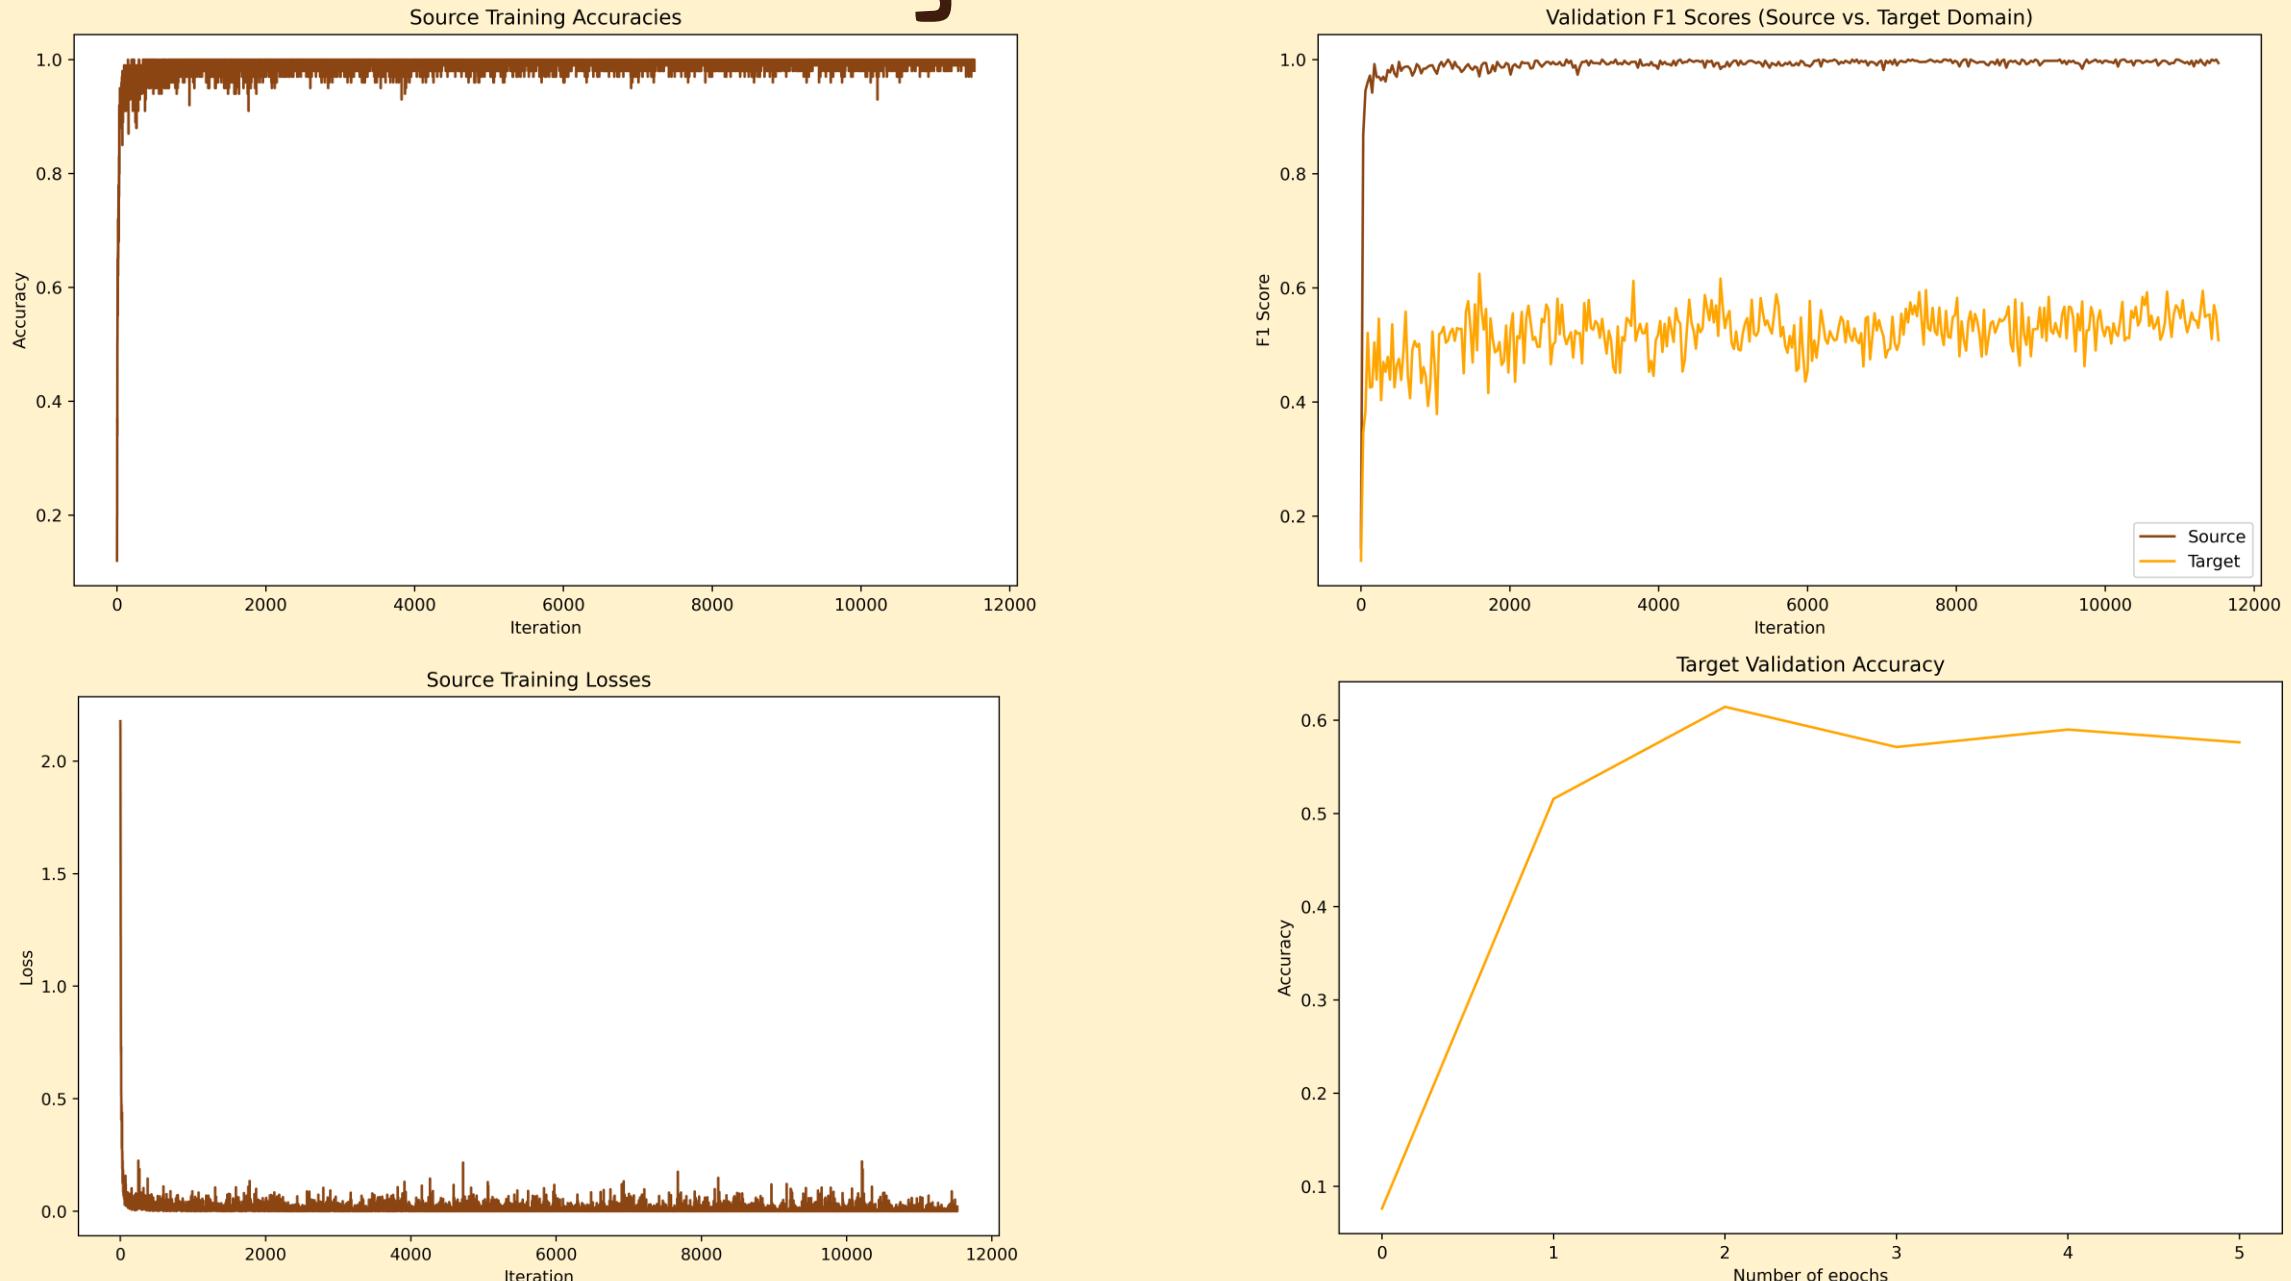
Gamma (Focal Loss)	Learning Rate	Lambda Max (DA factor)	Best Validation Accuracy
2	0.001	0.1	80.9%
2	0.001	1	83.2%
2	0.001	10	74.7%
2	0.0005	0.1	80.7%
2	0.0005	1	83.6%
2	0.0005	10	84.2%
5	0.001	0.1	81.1%
5	0.001	1	81.4%
5	0.001	10	83.5%
5	0.0005	0.1	82.5%
5	0.0005	1	85.4%
5	0.0005	10	84.4%

CORAL

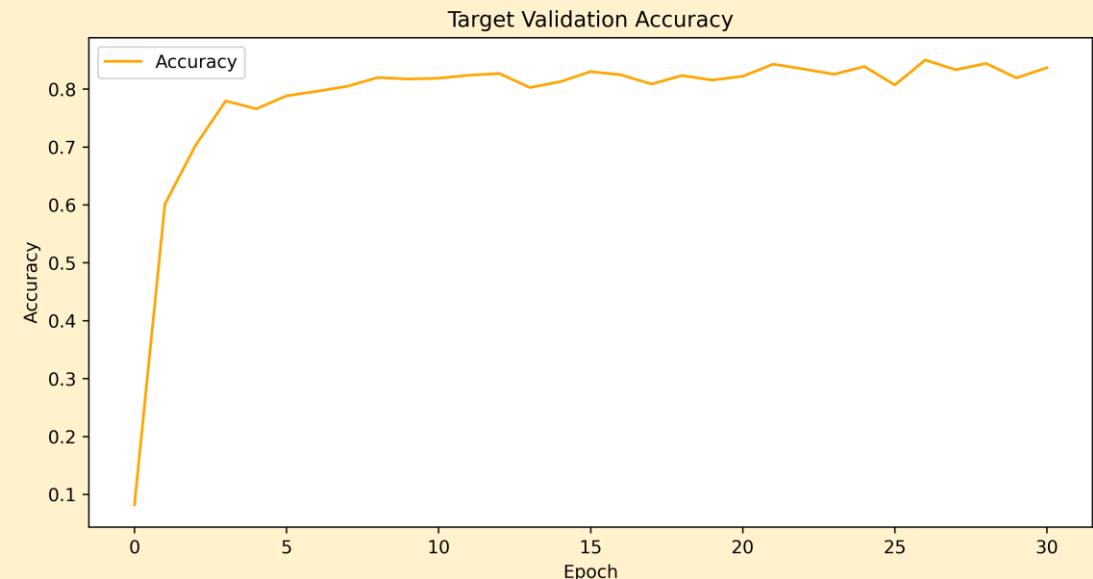
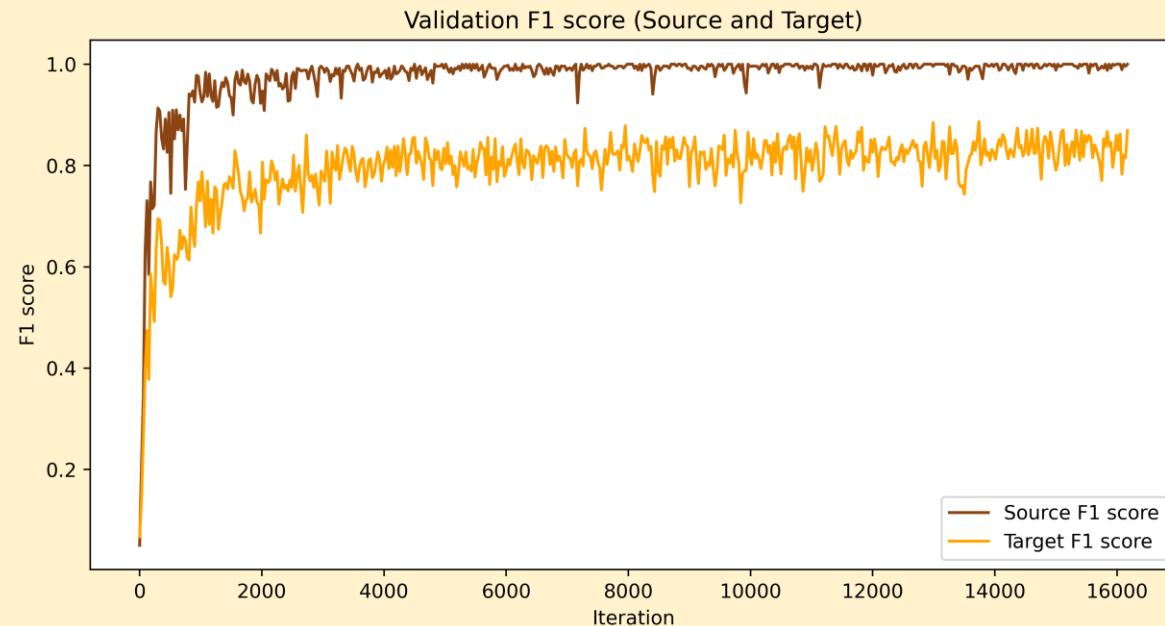
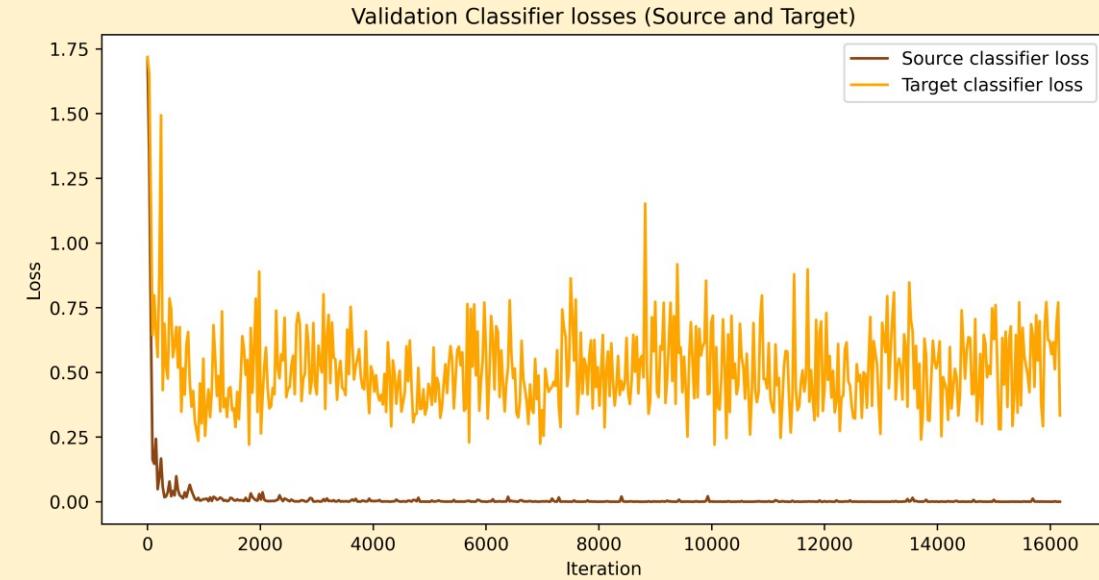
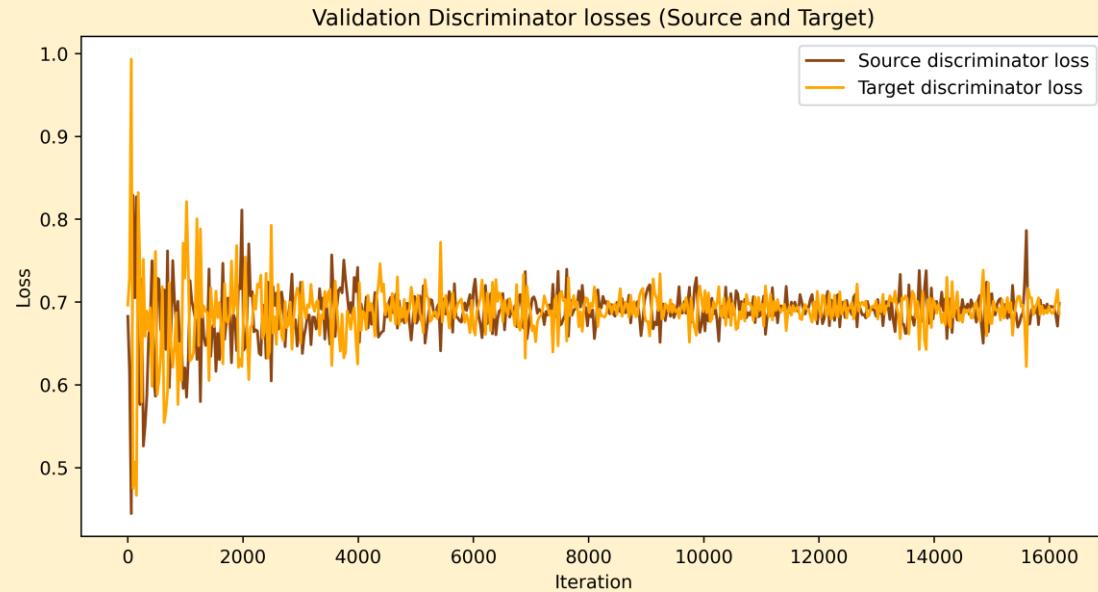
Gamma (Focal Loss)	Learning Rate	Lambda (DA factor)	Best Validation Accuracy
2	0.02	0.2	82.3%
2	0.02	variable	82.4%
2	0.02	0.3	84.4%
2	0.01	0.2	84.2%
2	0.01	variable	82.6%
2	0.01	0.3	78.8%
5	0.02	0.2	84.8%
5	0.02	variable	79.9%
5	0.02	0.3	85.0%
5	0.01	0.2	78.9%
5	0.01	variable	77.7%
5	0.01	0.3	76.5%

DANN

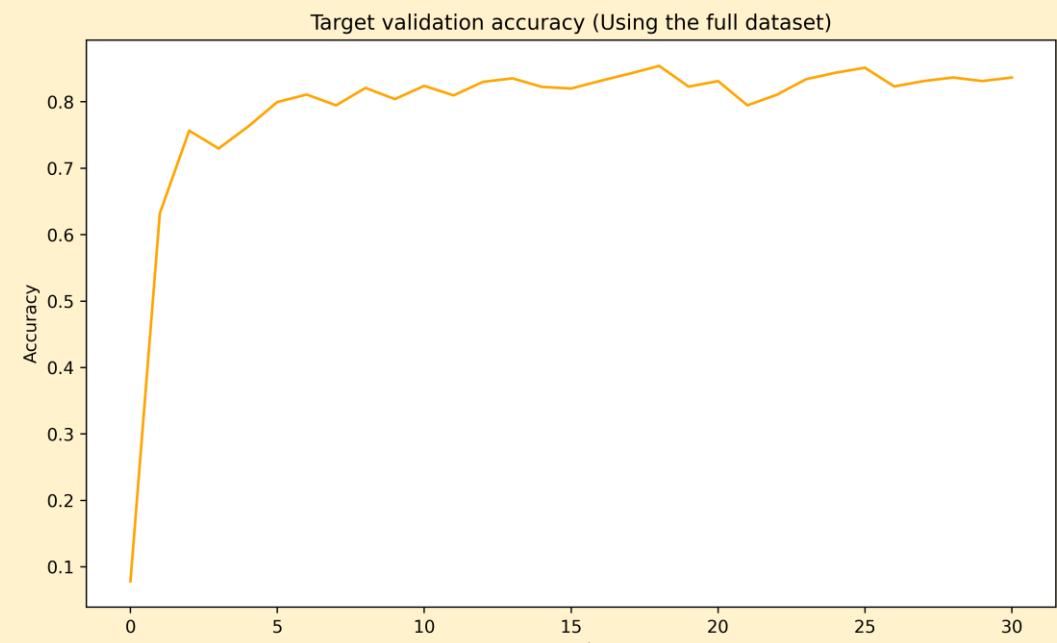
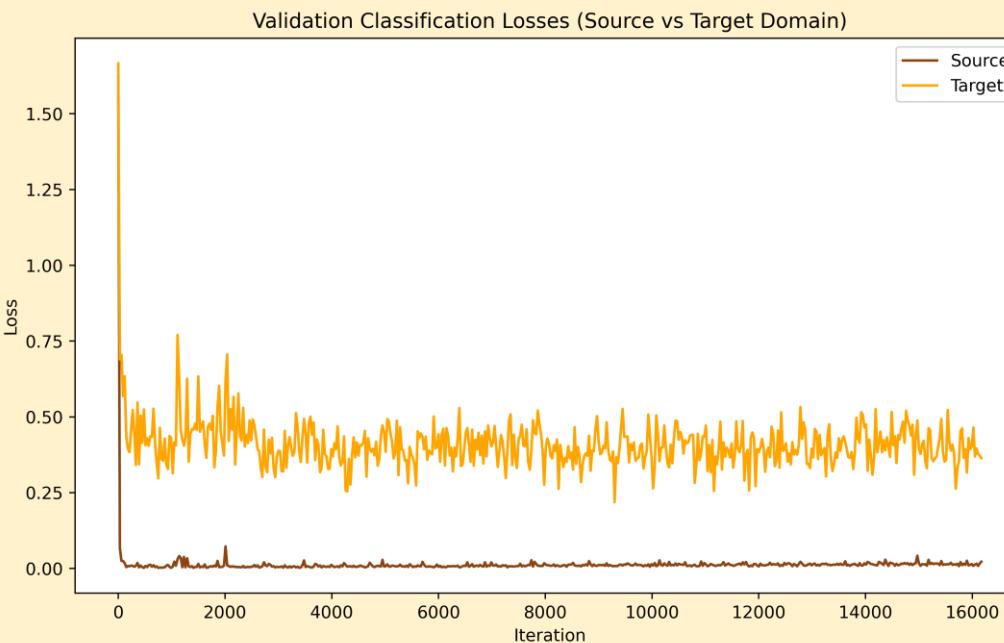
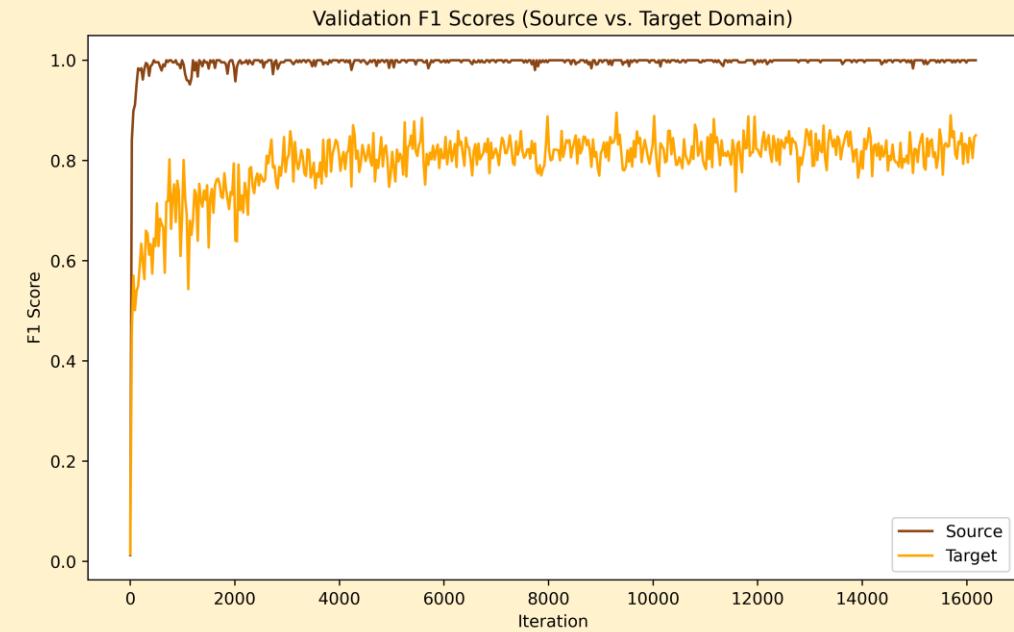
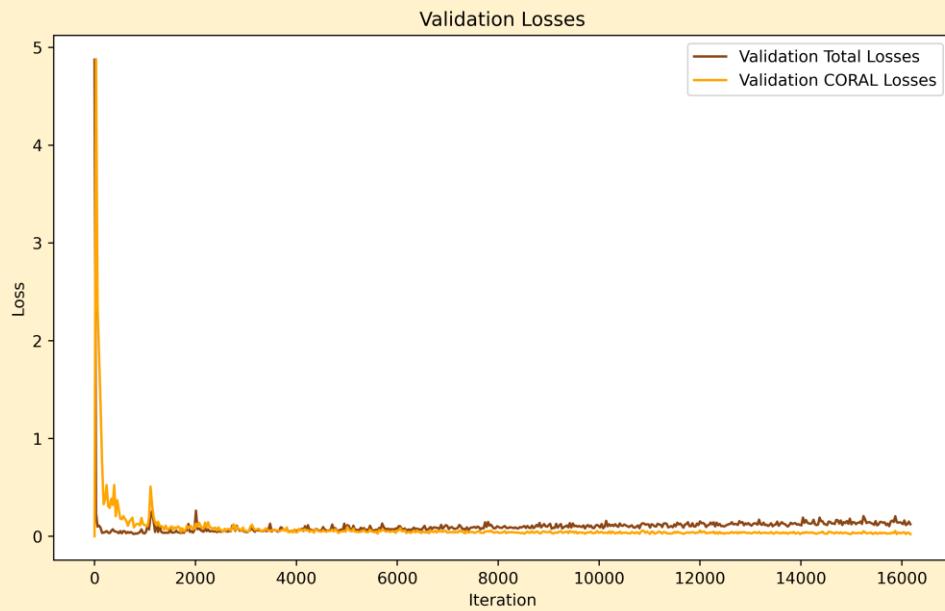
# Convergence: BASE-source



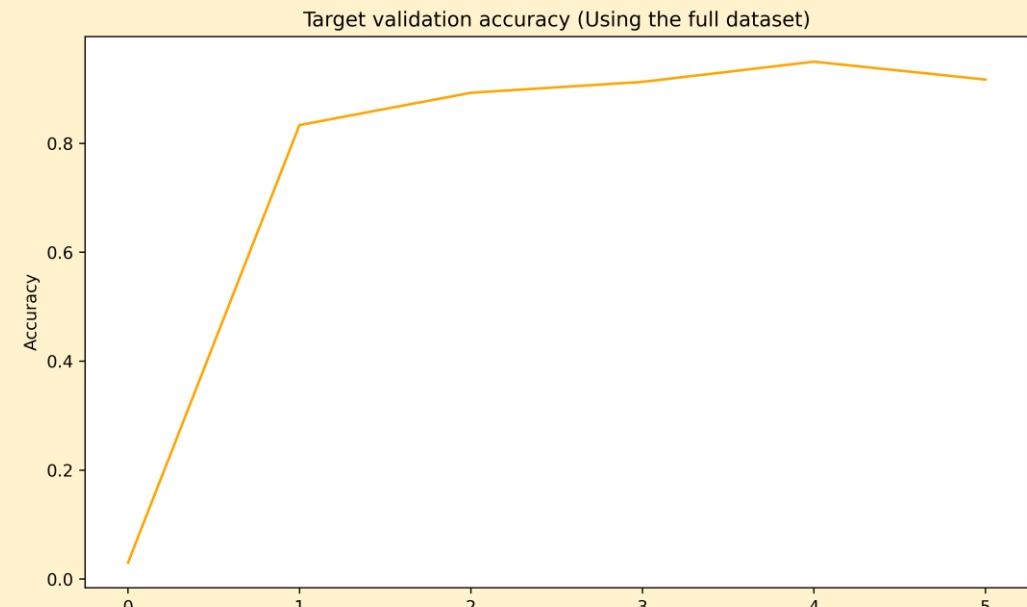
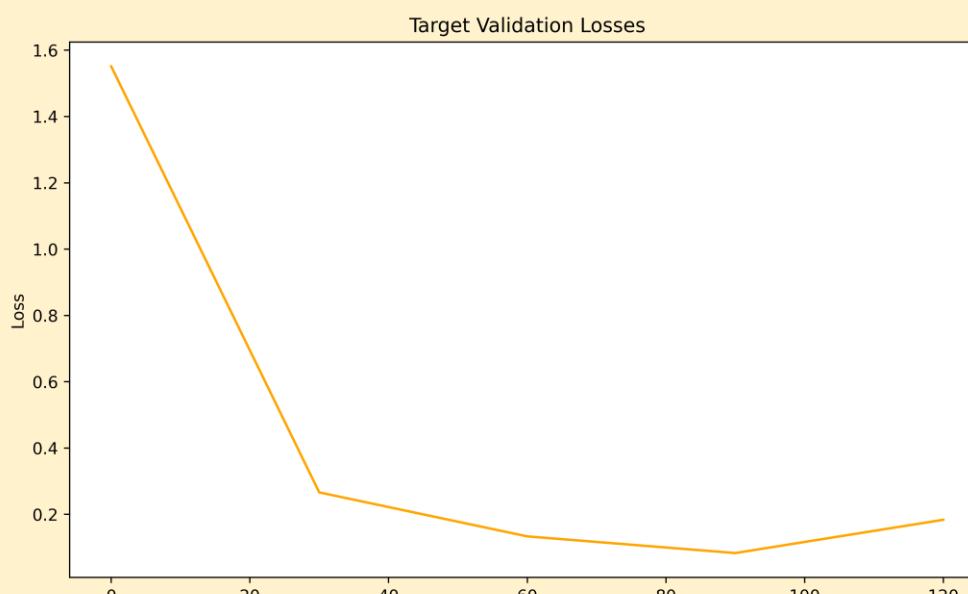
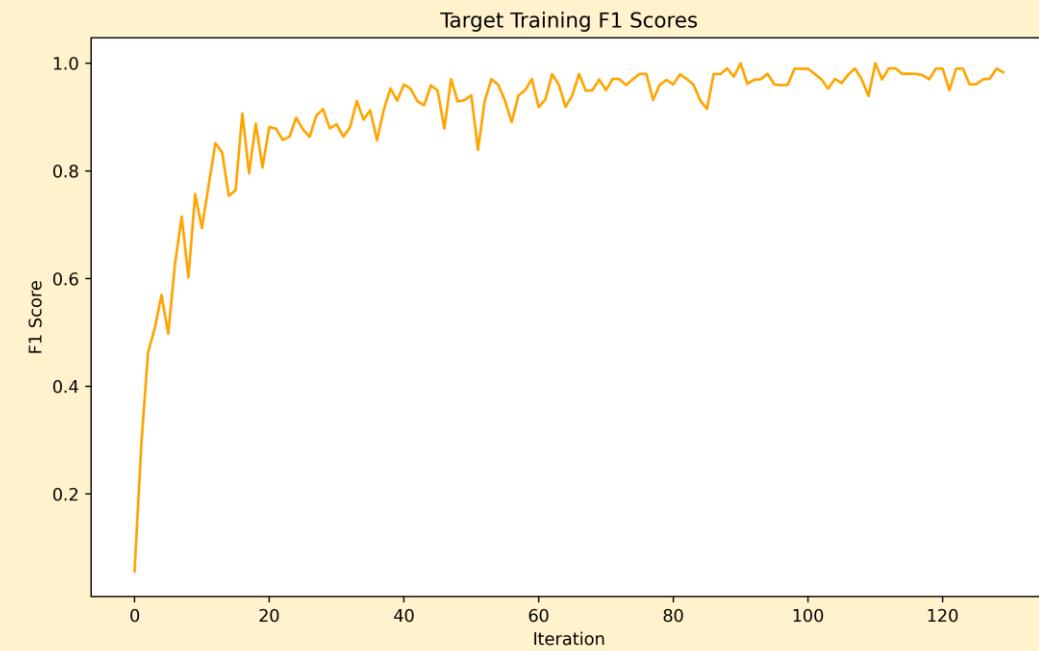
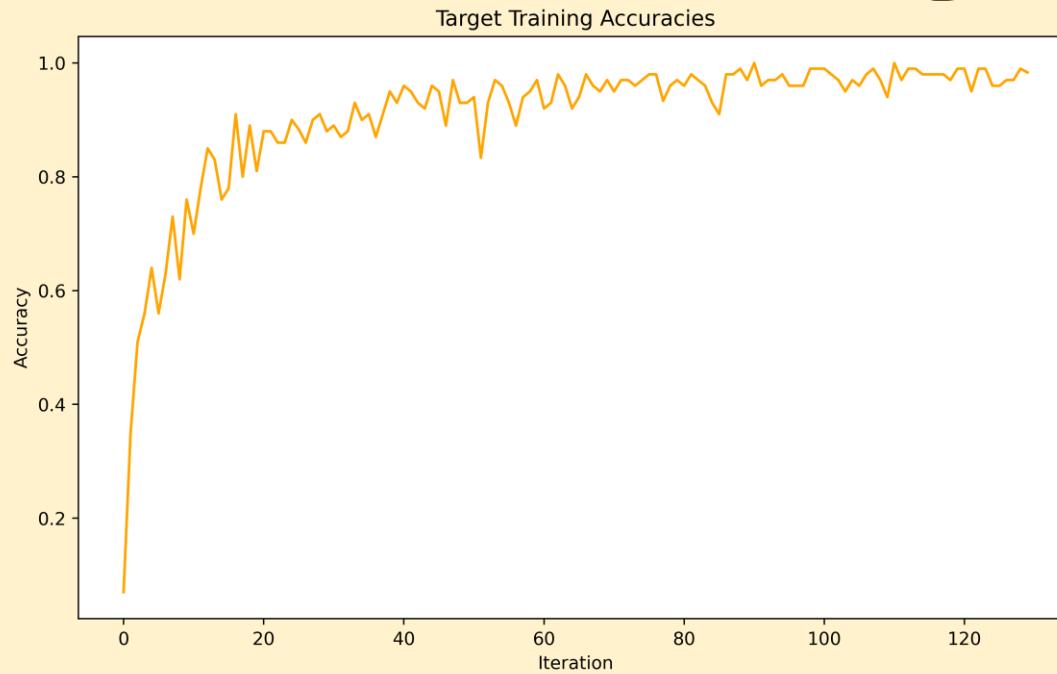
# Convergence: DANN



# Convergence: CORAL



# Convergence: BASE-target



# Results: Testing metrics on balanced source domain

Model	Testing Accuracy	F1-Score	Average AUPRC
BASE-source	99.17%	0.992	0.992
CORAL	99.94%	0.999	0.999
DANN	99.68%	0.997	0.997
BASE-target	36.56%	0.340	0.429

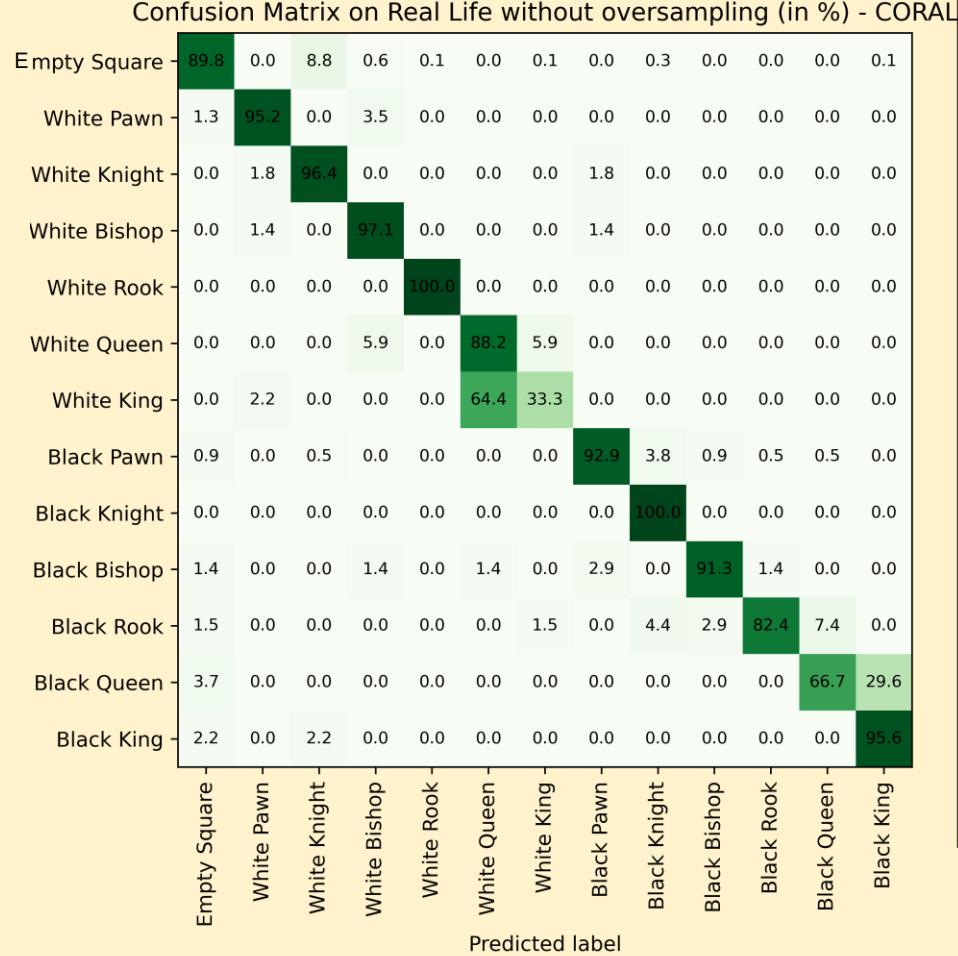
# Results: Testing metrics on balanced target domain

Model	Testing Accuracy	F1-Score	Average AUPRC
BASE-source	57.59%	0.553	0.595
CORAL	85.43%	0.852	0.859
DANN	83.59%	0.832	0.842
BASE-target	93.00%	0.930	0.937

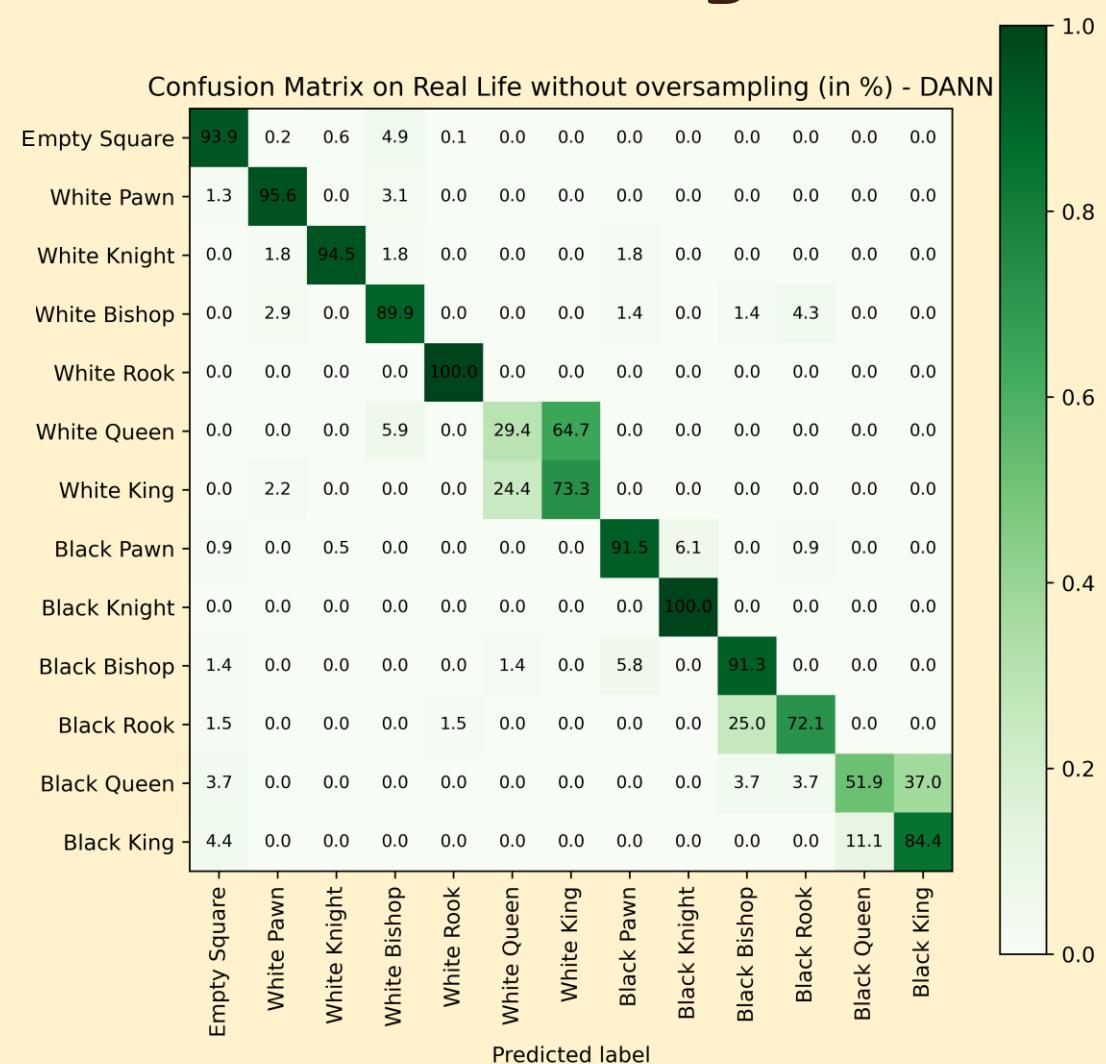
# Results: Testing metrics on imbalanced target domain

Model	Testing Accuracy	F1-Score	Average AUPRC
BASE-source	87.43%	0.874	0.556
CORAL	89.91%	0.919	0.825
DANN	92.31%	0.932	0.793
BASE-target	95.18%	0.956	0.889

# Results: Confusion matrices on imbalanced target domain

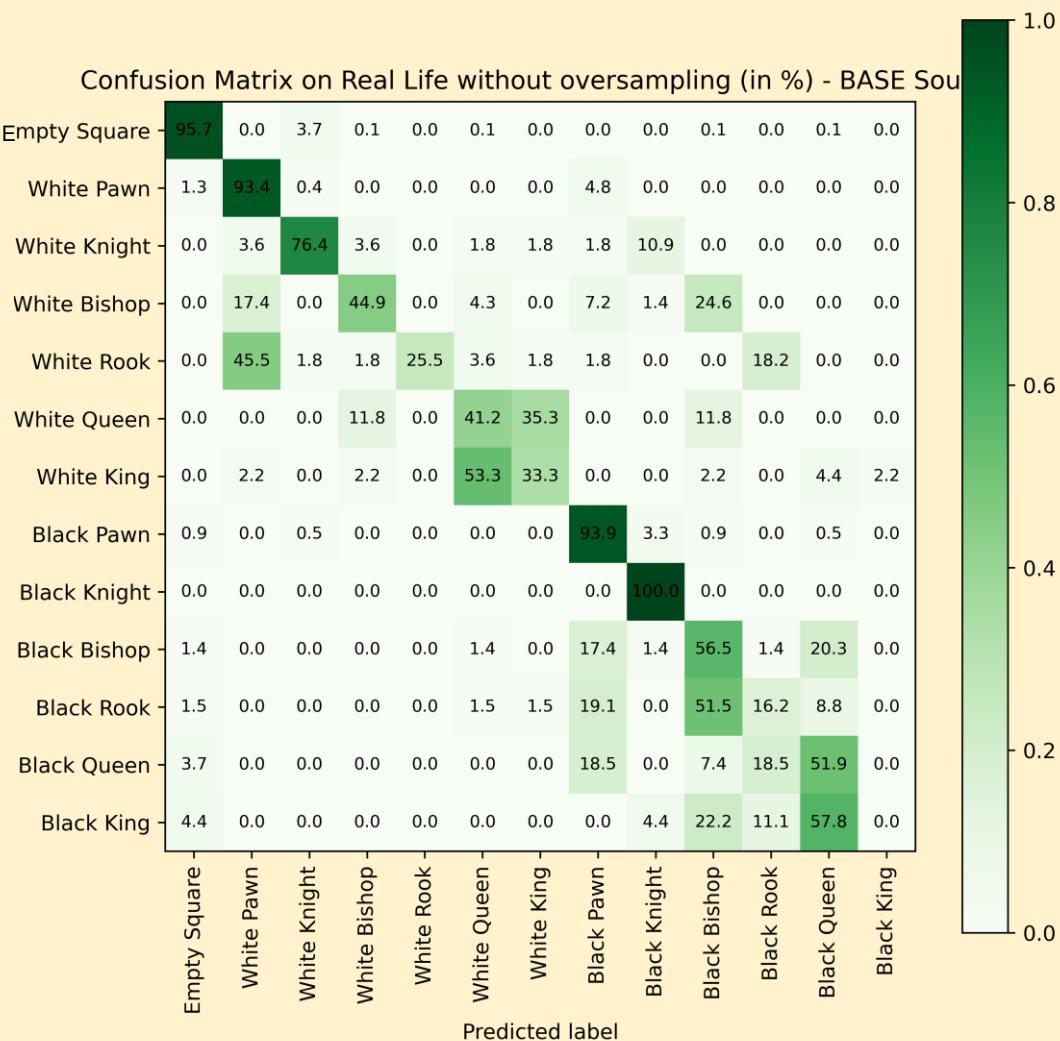


CORAL

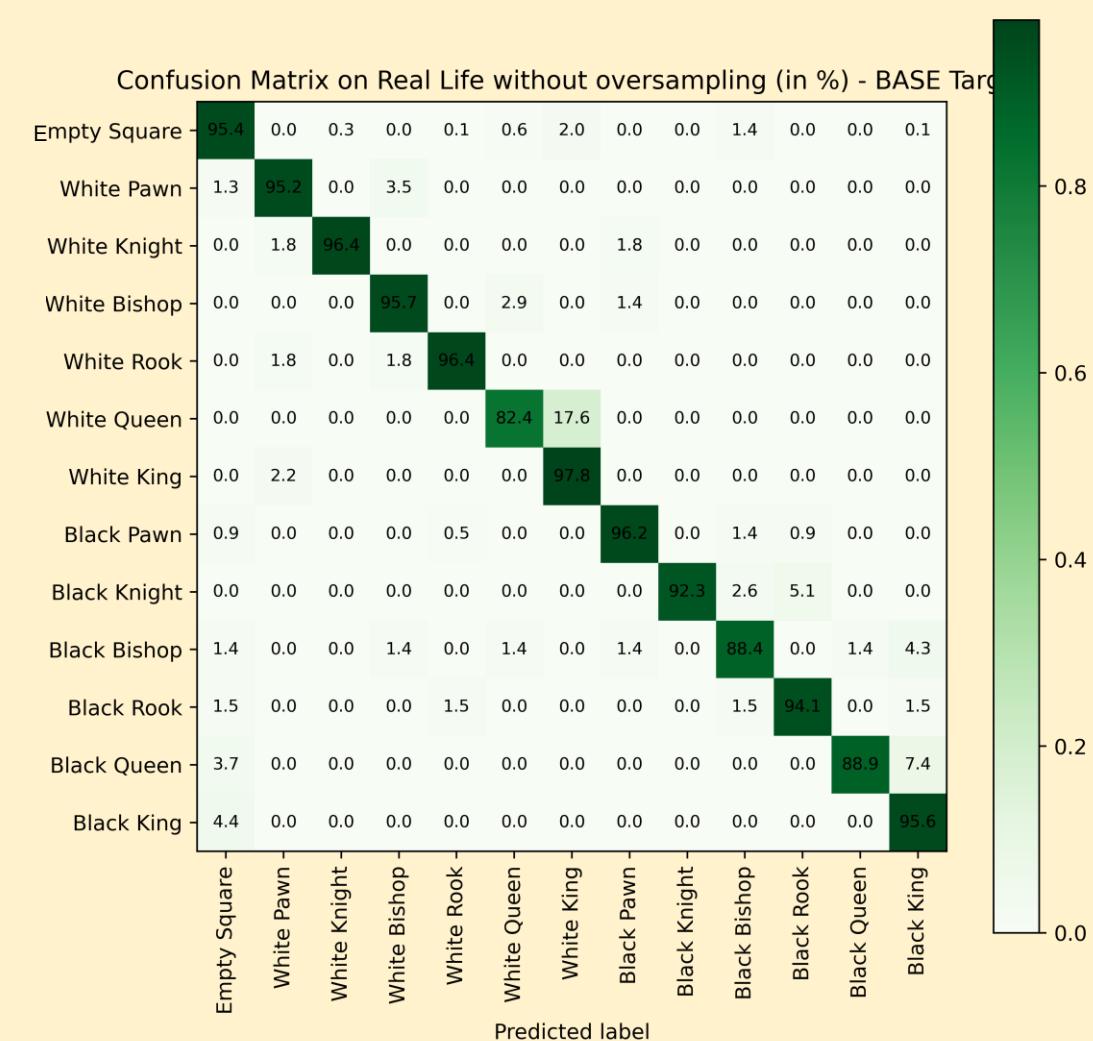


DANN

# Results: Confusion matrices on imbalanced target domain



BASE-source



BASE-target

# Conclusions

DANN & CORAL achieved results comparable to BASE-target

Acceptable results without requiring data labelling (For the most part)

Abundance of source domain data

Flexibility in manually tuning the feature & target distributions



Data generation was not straightforward at first and took time to R&D

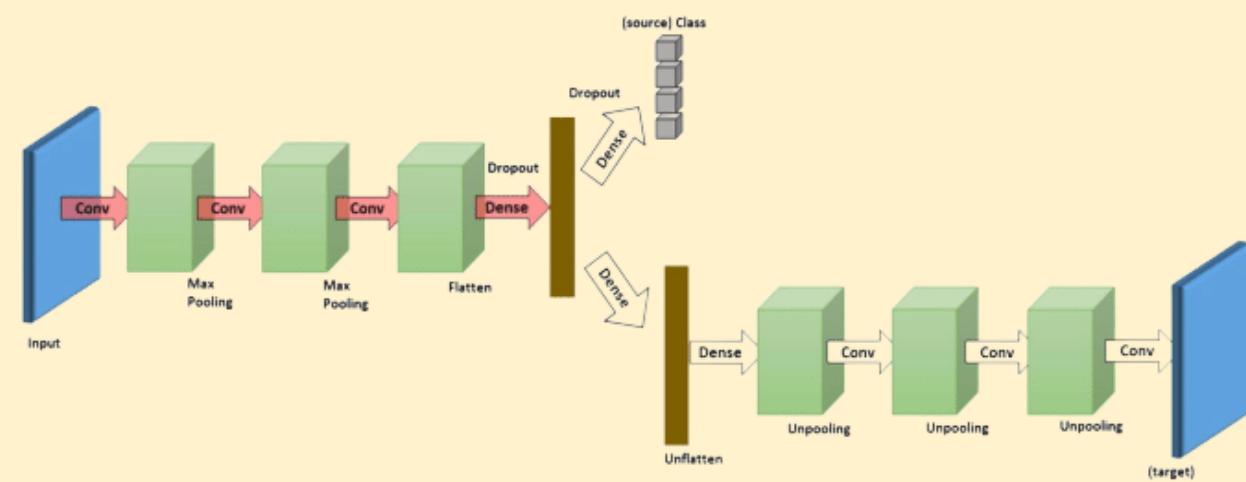
Hard time distinguishing kings & queens, a downside of using the top-view

Not invariant to piece set or board texture

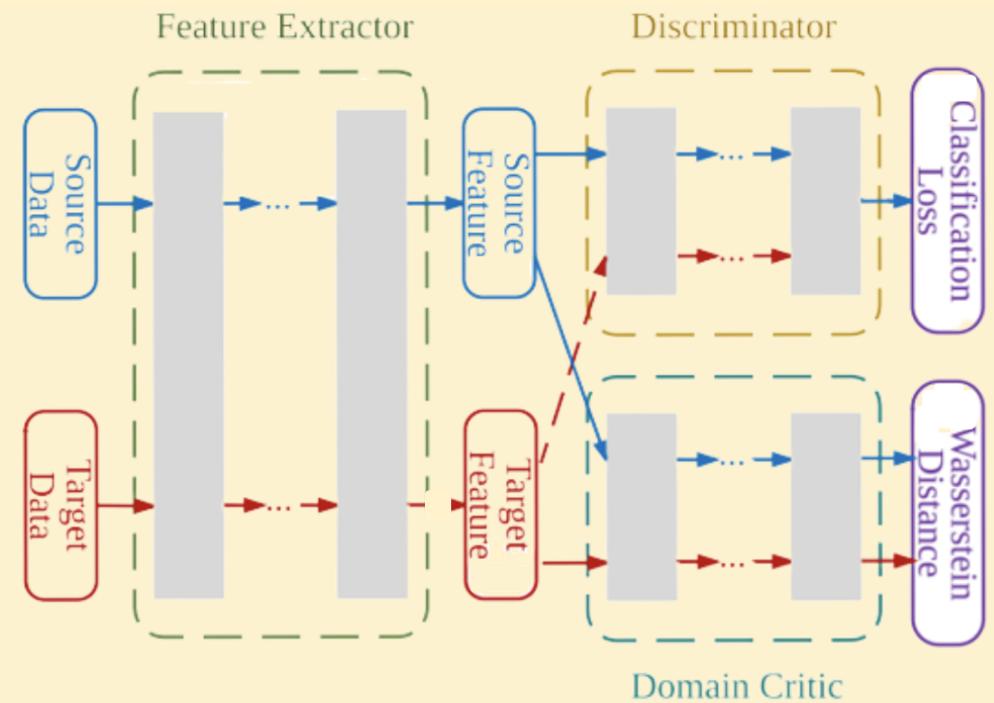
Does not utilize the rules of chess as natural prediction constraints

# Other possible approaches

## Alternative 1: Deep-Reconstruction Classification Network (DRCN)<sup>[10]</sup>



## Alternative 2: Wasserstein Distance Guided Representation Learning Model (WDGRL)<sup>[11]</sup>



[10]: M. Ghifary et al., "Deep reconstruction-classification networks for unsupervised domain adaptation," in Proc. Eur. Conf. Comput. Vis. (ECCV), pp. 597-613, 2016.

[11]: J. Shen et al., "Wasserstein distance guided representation learning for domain adaptation," in Proc. AAAI Conf. Artif. Intell., vol. 32, no. 1, 2018.

# References

- [1] S. Ben-David et al., "A theory of learning from different domains," *Mach. Learn.*, vol. 79, pp. 151-175, 2010.
- [2] A. D. S. D. Neto and R. M. Campello, "Chess position identification using pieces classification based on synthetic images generation and deep neural network fine-tuning," in Proc. 21st Symp. Virtual Augmented Reality (SVR), pp. 152-160, IEEE, Oct. 2019.
- [3] Elucidation, "Elucidation/chessboarddetect: Hodgepodge of chessboard detection algorithms on images from actual matches.", GitHub. [Online]. Available: <https://github.com/Elucidation/ChessboardDetect>. [Accessed: 10-Apr-2023].
- [4] "Fen to image," Chessvisionai RSS, <https://chessvision.ai/docs/tools/fen2image/> (accessed Jun. 1, 2023).
- [5] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, 'Focal Loss for Dense Object Detection', arXiv [cs.CV]. 2018.
- [6] B. Sun, J. Feng, and K. Saenko, 'Correlation Alignment for Unsupervised Domain Adaptation', CoRR, vol. abs/1612.01939, 2016.
- [7] K. Simonyan and A. Zisserman, 'Very Deep Convolutional Networks for Large-Scale Image Recognition', CoRR, vol. abs/1409.1556, 2014.
- [8] B. Sun, J. Feng, and K. Saenko, 'Correlation Alignment for Unsupervised Domain Adaptation', CoRR, vol. abs/1612.01939, 2016.
- [9] Y. Ganin et al., 'Domain-Adversarial Training of Neural Networks', arXiv [stat.ML]. 2016.
- [10]: M. Ghifary et al., "Deep reconstruction-classification networks for unsupervised domain adaptation," in Proc. Eur. Conf. Comput. Vis. (ECCV), pp. 597-613, 2016.
- [11]: J. Shen et al., "Wasserstein distance guided representation learning for domain adaptation," in Proc. AAAI Conf. Artif. Intell., vol. 32, no. 1, 2018.

Questions?